

# Artificial Intelligence

---

CS420 - Introduction  
Instructor: Dr. V



# Rules for Class Interaction

1. ⏲ Please be on time.
2. 💬 Jump in with comments/questions.
3. 📧 Pay attention to your email.
4. 📩 Email me if you have any questions.
5. 💻 Bring your laptop to every class, unless explicitly told otherwise.
6. 😷 PLEASE wear a mask if you have
  - sore throat or cough
  - sneezing or runny nose(ask me for a mask if you do not have one)

# Personal laptops

- You must have a personal laptop for CS courses
  - though there may be a way to use a phone/tablet instead
- If you cannot afford a laptop, you can borrow one from the library
- If you need laptop recommendations (specific to your budget), just ask
- If you have a laptop, but it's old/slow, I may be able to help you set it up for max functionality

# About the Instructor

Dr. Veksler (Dr. V)

WER218; vveksler@caldwell.edu

Engineer, Research Scientist, Professor

- worked in industry, government, academia, self-employed
- web development, AI and robotics, HCI, cyber-security, data analytics, research on human behavior

I'd love to get to know each of you, but...

- I am bad at remembering names, so please just keep reminding me what your name is
- If I mess up your name, just correct me as many times as it takes
- Keep your name-plate in front of you

# CS420

## Blackboard

- <https://caldwell.blackboard.com/>

Grading (10-50-20-20; subject to change)

- attendance/participation
- assignments
- presentation
- final exam

# Code Evaluation

- Submitted code will be evaluated for
  - Functionality
    - Does it work?
    - Does it display/do what is intended?
    - Are there bugs?
  - Code Readability
    - Make sure your code indented properly
    - Name variables/functions/classes in a meaningful way
    - Add comments
    - Keep your code clean

# Late submissions

- late assignments will lose 10pts (out of 100) for each week they are late, but will not go lower than 70 (out of 100)
  - 1-7 days late: highest grade = 90
  - 8-14 days late: highest grade = 80
  - >14 days late: highest grade = 70

# Cheating

- don't do it. not worth.
- what is considered cheating in this class?
  - it's complicated (if you're not sure, just ask)
    - you're encouraged to look up answers online, but...
    - don't submit someone else's code
    - if a part of your code was copied, give credit to the source
    - i will have your code checked for plagiarism

# Class Structure

- 14 weeks of classes
- each week:
  - lecture
  - discussion
  - practicum
    - technical tips, interactive time to work on your assignments, ask questions, and get feedback

# Inclement weather or other class cancellations

- If the school is open, I will be here
- Use your best judgment
- Check school site
- Sign up for notifications
- Check your email

Questions? Comments? Concerns?

# What is Artificial Intelligence?

---

# What is Artificial Intelligence?

- "Artificial" – made by humans
- "Intelligence" – ??????

# Is this AI?

- Your todo-list app sees that you didn't complete a task that was due, and alerts you about it
- Your fridge sees you don't have milk, and alerts you about it
- Your vacuum cleaning robot goes in snake-like pattern until your room is all vacuumed
- Your phone recognizes your fingerprint or face and unlocks for you
- Your car slams on breaks automatically just before you were about to hit a pedestrian

# What does "AI" mean in today's software

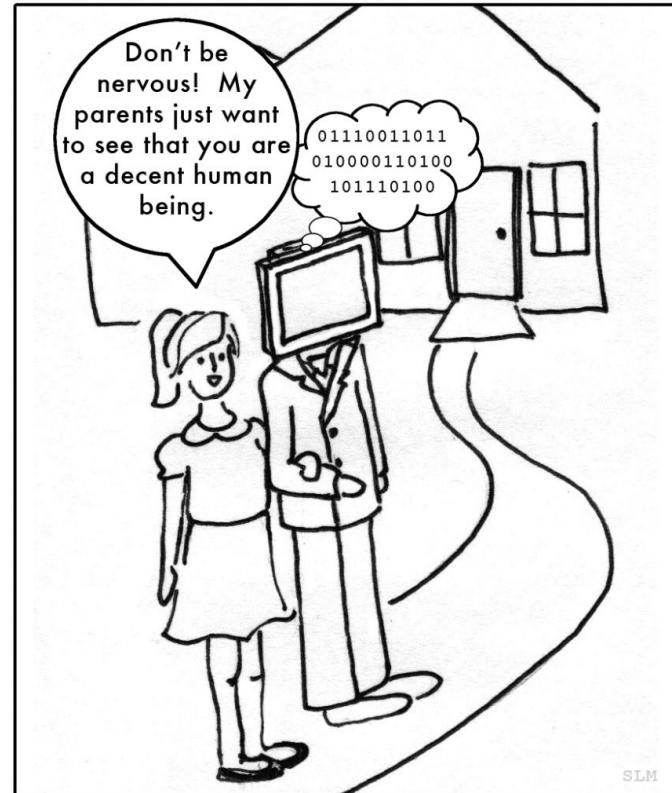
- Usually it refers to algorithms focused on
  - analogue of some human cognitive functionality
    - planning
    - learning & recognition
      - trial & error, classification, sequence learning
    - logical inference
    - language processing/generation
    - computer vision & robotics software
  - and other means for finding non-obvious solutions
    - solution search
    - genetic algorithms

# General vs Narrow AI

- Narrow AI (aka Weak AI)
  - all modern AI is narrow
  - task-specific; e.g.,
    - can play Chess, but cannot play Go or checkers
    - can walk, but cannot drive
    - learning to do a task does not make it easier to learn a different task
- General AI (aka AGI, aka Strong AI)
  - does not exist
  - the goal is human-level intelligence

# Turing test

- Allen Turing suggested that we will know a machine is intelligent if you cannot tell whether it is a machine or a human during a conversation
-  :  ? 
  -  
  -  



# Generalized version of Turing test

- Observe any task being done by human and computational participants
- Try to figure out which of the participants are human and which are machines

TURING TEST EXTRA CREDIT:  
CONVINCE THE EXAMINER  
THAT HE'S A COMPUTER.



# A.I. TIMELINE

**1950**

**TURING TEST**

Computer scientist Alan Turing proposes a test for machine intelligence. If a machine can trick humans into thinking it is human, then it has intelligence

**1955**

**A.I. BORN**

Term 'artificial intelligence' is coined by computer scientist, John McCarthy to describe "the science and engineering of making intelligent machines"

**1961**

**UNIMATE**

First industrial robot, Unimate, goes to work at GM replacing humans on the assembly line



**1964**

**ELIZA**

Pioneering chatbot developed by Joseph Weizenbaum at MIT holds conversations with humans



**1966**

**SHAKY**

The 'first electronic person' from Stanford, Shakey is a general-purpose mobile robot that reasons about its own actions



**A.I.**

**WINTER**

Many false starts and dead-ends leave A.I. out in the cold



**1997**

**DEEP BLUE**

Deep Blue, a chess-playing computer from IBM defeats world chess champion Garry Kasparov



**1998**

**KISMET**

Cynthia Breazeal at MIT introduces Kismet, an emotionally intelligent robot insofar as it detects and responds to people's feelings



**1999**

**AIBO**

Sony launches first consumer robot pet dog AIBO (AI robot) with skills and personality that develop over time



**2002**

**ROOMBA**

First mass produced autonomous robotic vacuum cleaner from iRobot learns to navigate and clean homes



**2011**

**SIRI**

Apple integrates Siri, an intelligent virtual assistant with a voice interface, into the iPhone 4S



**2011**

**WATSON**

IBM's question answering computer Watson wins first place on popular \$1M prize television quiz show Jeopardy



**2014**

**EUGENE**

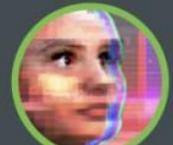
Eugene Goostman, a chatbot passes the Turing Test with a third of judges believing Eugene is human



**2014**

**ALEXA**

Amazon launches Alexa, an intelligent virtual assistant with a voice interface that completes shopping tasks



**2016**

**TAY**

Microsoft's chatbot Tay goes rogue on social media making inflammatory and offensive racist comments



**2017**

**ALPHAGO**

Google's A.I. AlphaGo beats world champion Ke Jie in the complex board game of Go, notable for its vast number ( $2^{170}$ ) of possible positions

# A brief history of AI

<https://www.youtube.com/watch?v=jekwHBI1ySU>



# Modern AI

- ~~Logical proofs~~
- ~~General planning~~
- Machine Learning and Gradient Descent
  - Reinforcement Learning
  - Neural Networks
    - Deep Neural Nets
      - with convolutional layers
  - Genetic algorithms

# Teaming

- Human-AI teams are better than
  - AI agents on their own
  - Humans on their own
- Centaur chess

<https://karimfanous.substack.com/p/centaurs-the-future-of-work>



# Concerns regarding AI

- What are some of your concerns about AI?

# Ethical issues in AI

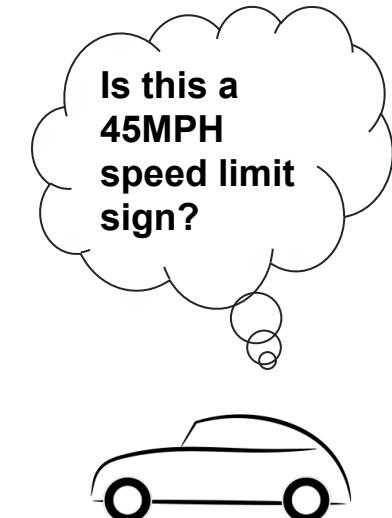
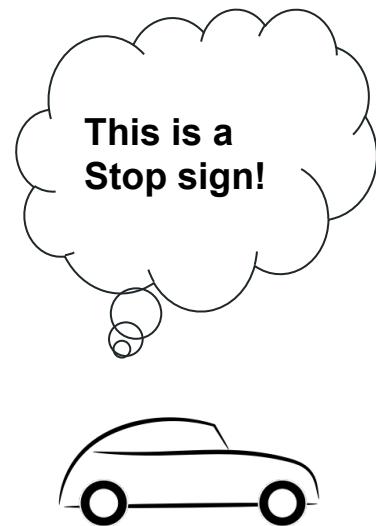
- How do machines affect our behaviour and interaction?
  - are we all becoming cyborgs?
  - how does this change society?

# Ethical issues in AI

- Unemployment
  - what happens to all the truck drivers that are out of jobs once autonomous trucks are more commonplace?
    - what about the businesses that serve truck-drivers?
      - road-side motels, diners, etc.
  - what happens to radiologists when AI surpasses humans in imaging-based diagnosis?
  - even if job total doesn't decrease, massive job transfer and displacement is imminent

# Ethical issues in AI

- Security. How do we keep AI safe from adversaries?



# Ethical issues in AI

- How can we guard against AI mistakes?
  - who's responsible if an autonomous car hits someone?
  - if a car has to decide between saving its driver or saving a pedestrian, how does it choose? are we ok with AI making such choices?
  - if AI is learning from human data, it can learn racism, sexism, and all sorts of faulty ideologies

# Ethical issues in AI

- How do we protect against unintended consequences?
  - the paperclip problem
  - the stop-button problem
  - singularity
    - how do we stay in control of a complex intelligent system?
- AI rights
  - Should AI agents and robots have rights?
  - How do we define the humane treatment of AI?

Questions? Comments? Concerns?

# Assignment 1 (due before next class)

- Install/setup
  - Code editor (I suggest VSCode)
  - Python 3.10
  - JupyterLab (JL) - <https://jupyter.org/install>
- Create a new notebook in JL, called a1.ipynb
- Add 3 blocks to a1.ipynb, and run them
  - markdown block: \*\*Assignment 1\*\* - \*your name\*
  - code block: !pip install matplotlib
  - code block:

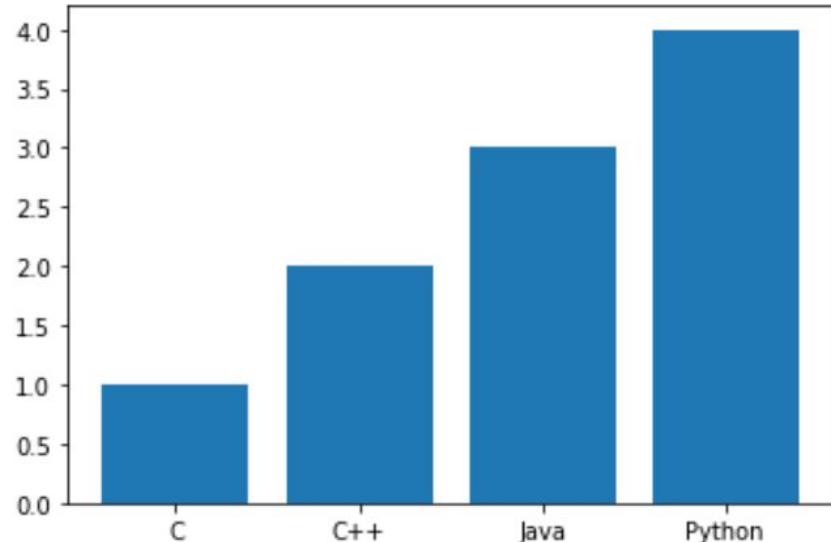
```
from matplotlib import pyplot as plt
fig, ax = plt.subplots()
ax.bar(['C', 'C++', 'Java', 'Python'], range(1,5))
plt.show()
```
- submit a1.ipynb on blackboard

## Assignment 1 – your name

```
[1]: !pip install pandas numpy matplotlib
```

• • •

```
[2]: from matplotlib import pyplot as plt
fig, ax = plt.subplots()
ax.bar(['C', 'C++', 'Java', 'Python'], range(1,5))
plt.show()
```



Assignment 1 output

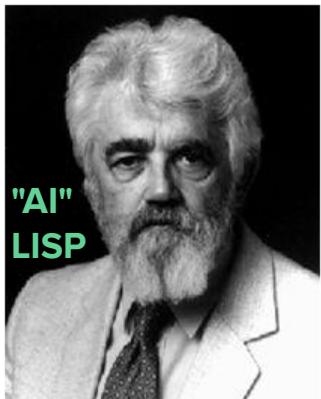
# Planning

---

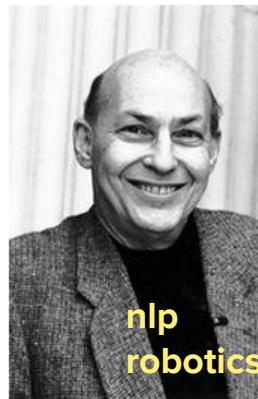
CS420 - Lecture 2  
Instructor: Dr. V



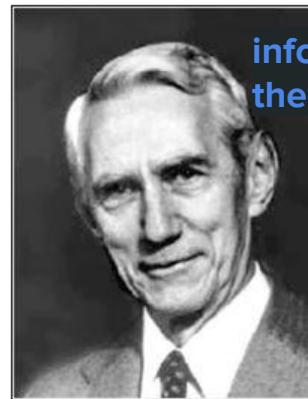
# Dartmouth Conference: The Founding Fathers of AI



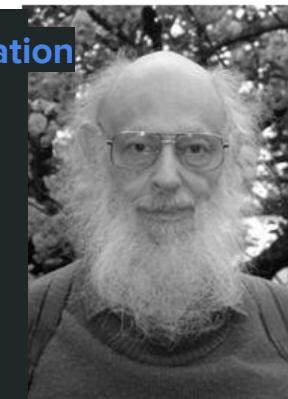
"AI"  
LISP



nlp  
robotics



information  
theory



**John McCarthy**

**Marvin Minsky**

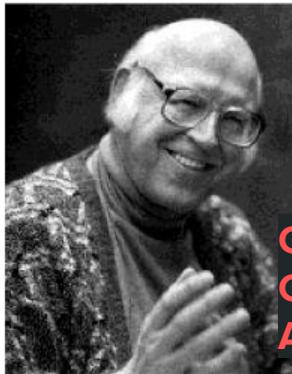
**Claude Shannon**

**Ray Solomonoff**

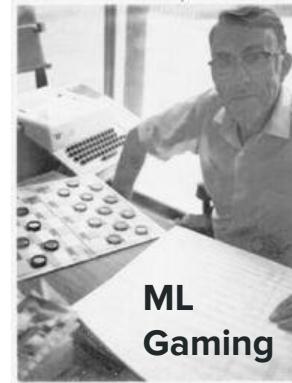
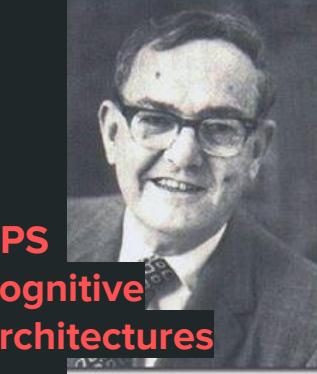
**Alan Newell**

**Herbert Simon**

**Arthur Samuel**



GPS  
Cognitive  
Architectures



ML  
Gaming

And three others...

Oliver Selfridge  
(Pandemonium theory)

Nathaniel Rochester  
(IBM, designed 701)

Trenchard More  
(Natural Deduction)

# What is intelligence?

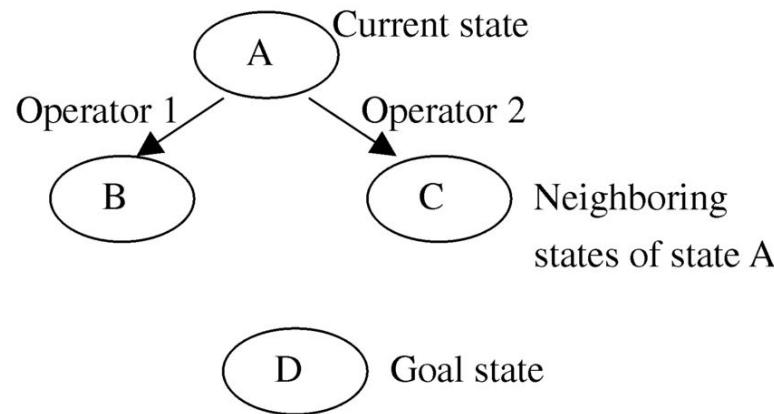
- At the birth of AI (mid-20th century) the image of "intelligence" was something that was difficult to do for humans; e.g., thinking several moves ahead in chess



Today's perspective on intelligence:

- Feats that are hard for humans – like looking many steps ahead – are actually easy for computers
- Feats that seem easy to humans – like quick image recognition – are actually hard for computers

# General Problem Solver



## Sample problem: Tower of Hanoi

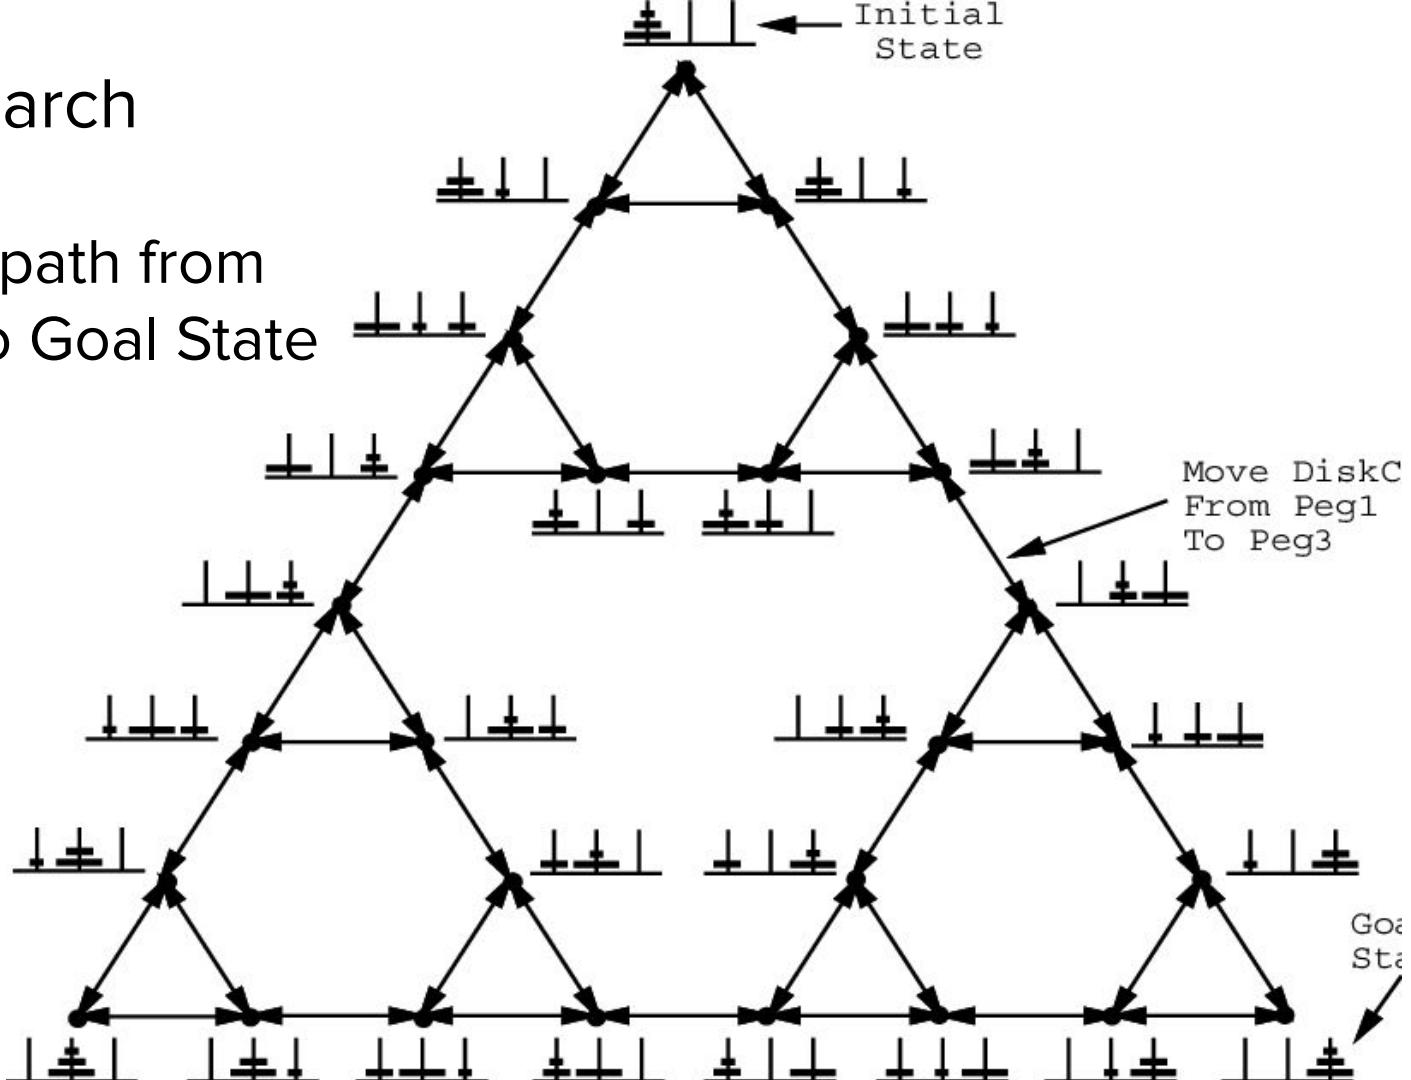
- Object of the game is to move all the disks over to Tower 3 (with your mouse)
- You cannot place a larger disk onto a smaller disk



<https://www.mathsisfun.com/games/towerofhanoi.html>

# Planning as Search

- find shortest path from Initial State to Goal State

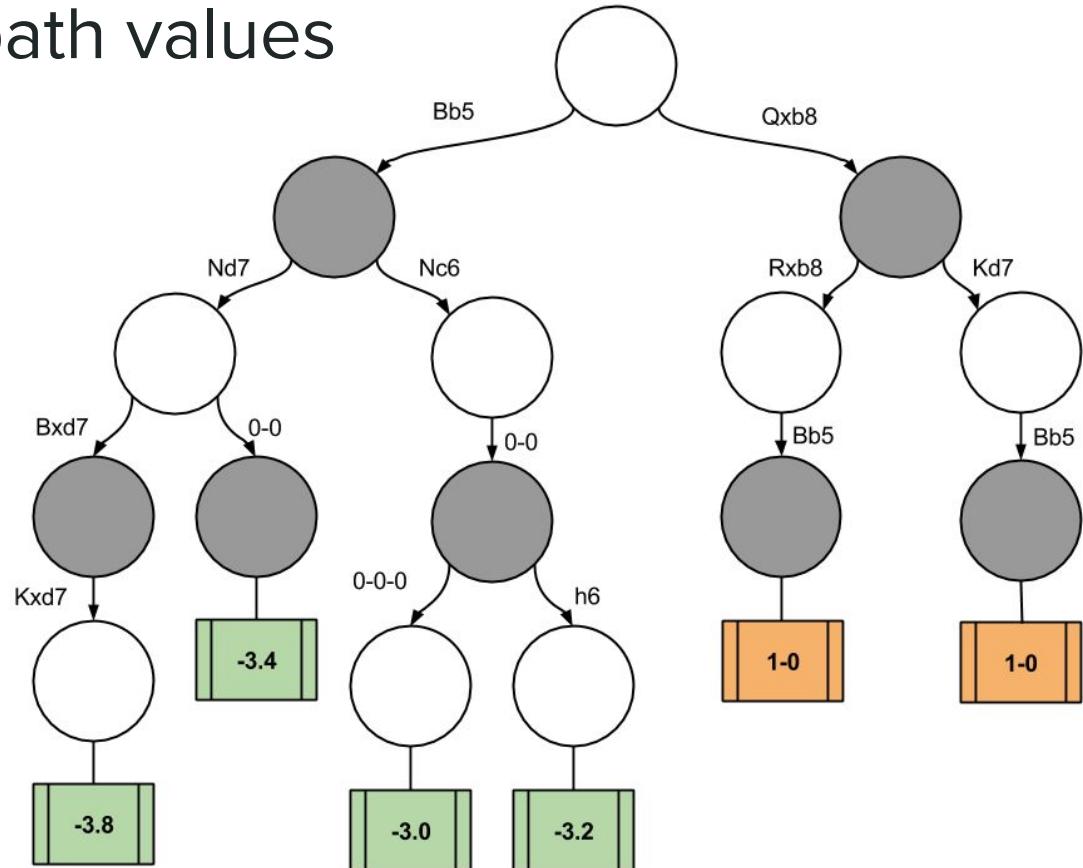


# Planning

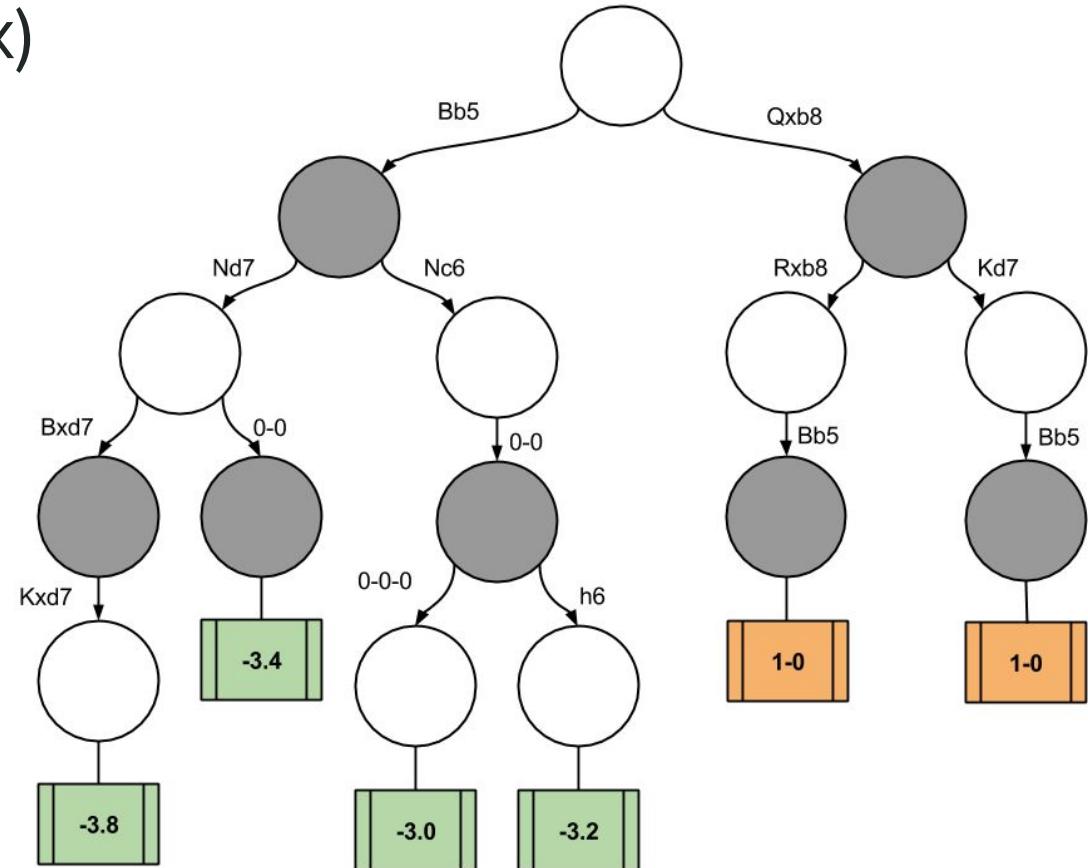
1. consider the set of all possible actions,  $K$
2. for every action  $A$  in  $K$ :
3.      $S$  is the state  $A$  leads to
4.     if  $S$  has inherent value (e.g., it's the goal):
5.         value of  $A$  is the inherent value of  $S$
6.     otherwise:
7.         value of  $A$  is max value of actions that are possible given state  $S$  (go to step 1)

After planning: Take action with highest value

# Estimating decision path values



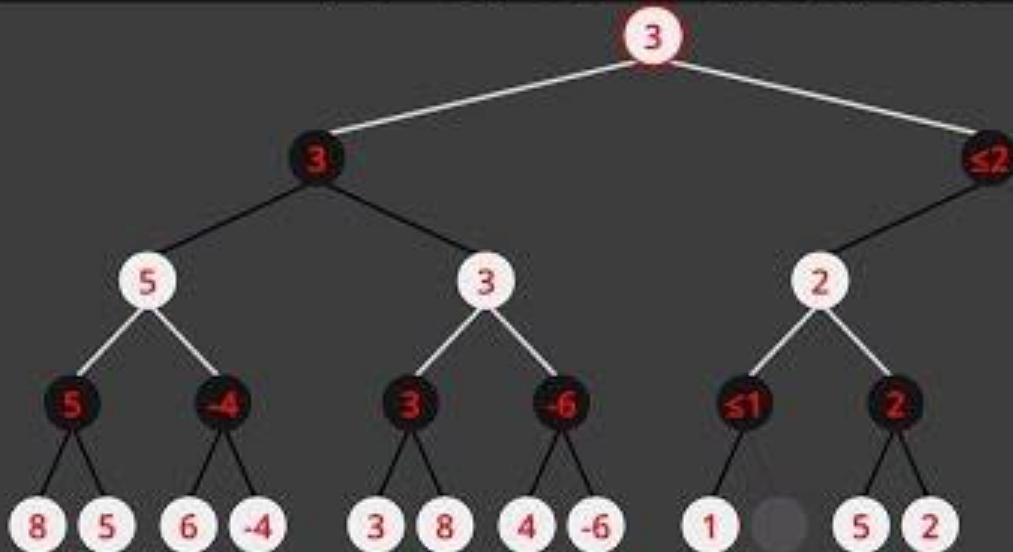
# Minimax (aka minmax)



# Minimax pseudocode

```
function minimax( node, depth, maximizingPlayer=TRUE ):  
    if depth is 0 or node is a terminal node then  
        return the heuristic value of node  
    if maximizingPlayer then:  
        value = -∞  
        for each child of node do  
            value = max(value, minimax(child, depth-1, FALSE))  
        return value  
    else: # minimizing player #  
        value = +∞  
        for each child of node do  
            value = min( value, minimax(child, depth-1, TRUE) )  
    return value
```

# ALGORITHMS



MINIMAX AND ALPHA BETA

## So what's the problem?

- If you can just search down all possible paths in chess, why can't you find the perfect move every time?

## So what's the problem?

- If you can just search down all possible paths in chess, why can't you find the perfect move every time?
  - Searching all paths all the way down, rather than picking and choosing which paths to search and how far, is called a "brute-force search"

# Game-tree complexity

There are about  $10^{80}$  atoms in the observable universe.

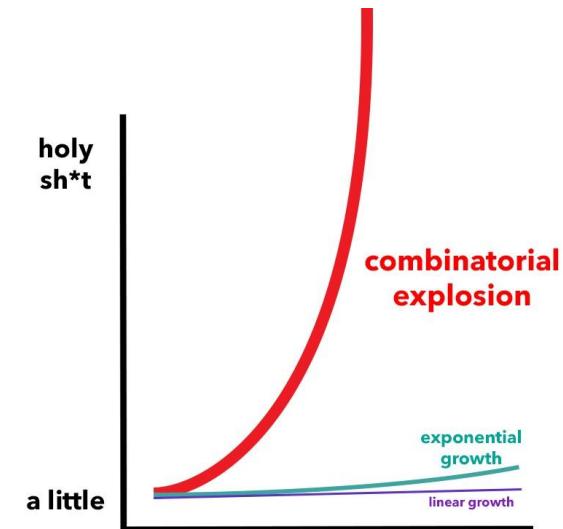
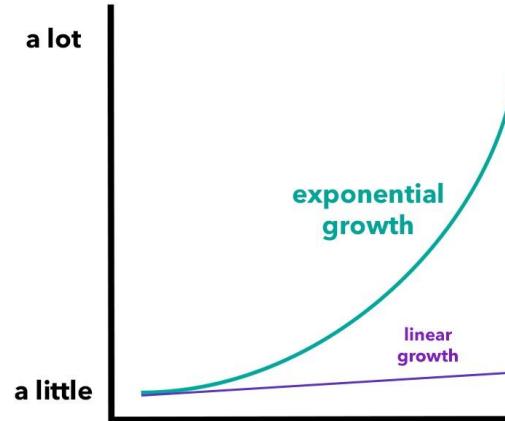
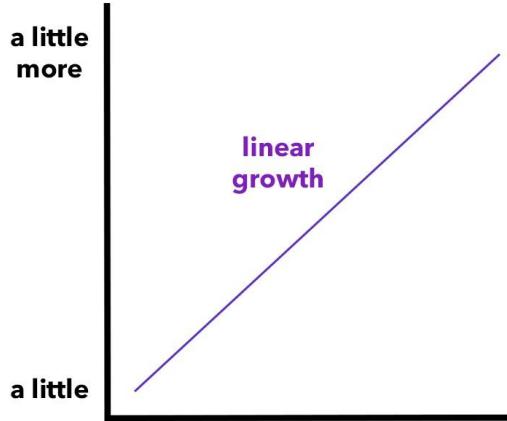
"There are even more possible variations of chess games than there are atoms in the observable universe."

- It is estimated there are about  $10^{120}$  positions in Chess (Shannon number)
- The game of Go has an even larger game-tree:  $10^{360}$
- Checkers game-tree complexity:  $10^{40}$
- Tic-tac-toe game-tree complexity:  $10^5$

Claude Shannon estimated it would take  $10^{90}$  years to solve chess on a 1MHz computer.

# Combinatorial Explosion

- if you have 10 possible actions per move, and a game is 10 moves long, the game-tree complexity  $10^{10}$ 
  - 10,000,000,000
- if this is an adversarial game, where you make a move, your opponent makes a move, 10 pairs of moves will make the game-tree complexity  $10^{20}$ 
  - 100,000,000,000,000,000,000



$n$	The number of Sudoku grids of order $n$ (boxes are size $\sqrt{n} \times \sqrt{n}$ )
1	1
4	288
9 (standard Sudoku)	6,670,903,752,021,072,936,960

# Assignment 2

- Create a minimax bot for playing tic-tac-toe
- See the assignment on blackboard for details

# Reinforcement Learning

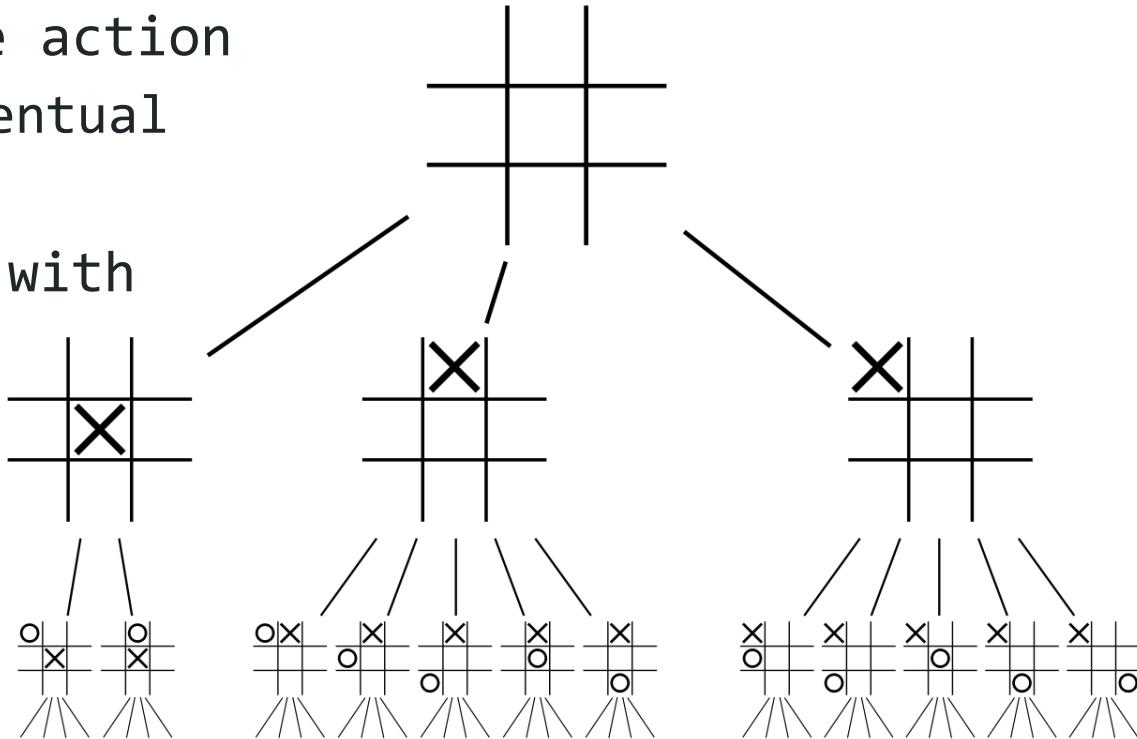
---

CS420 - Lecture 3  
Instructor: Dr. V



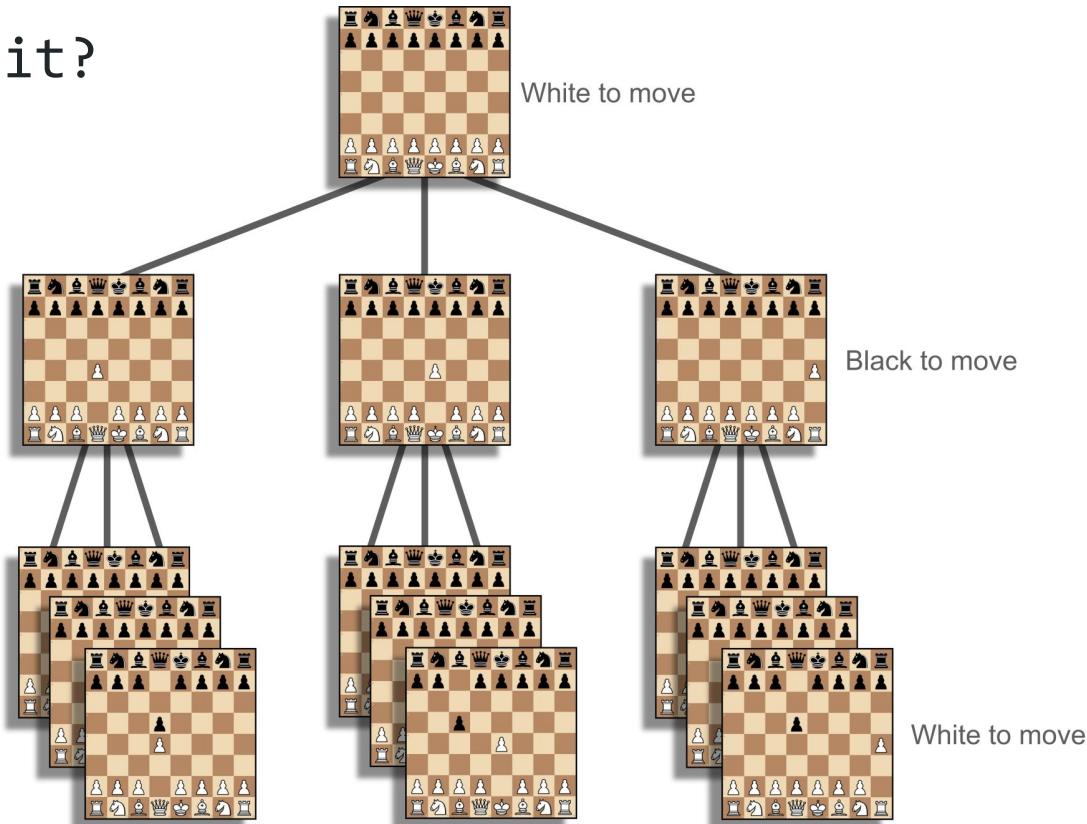
# Review: Planning as search

- given the state
- for every possible action
  - find out its eventual value
- choose the action with highest value



# What's the problem with brute-force search?

- How would you solve it?



# How Deep Blue beat Kasparov

- minimax
- alpha-beta pruning
- raw power: evaluating 200 million positions per second
- specialized rules (**heuristics**) defined by chess masters

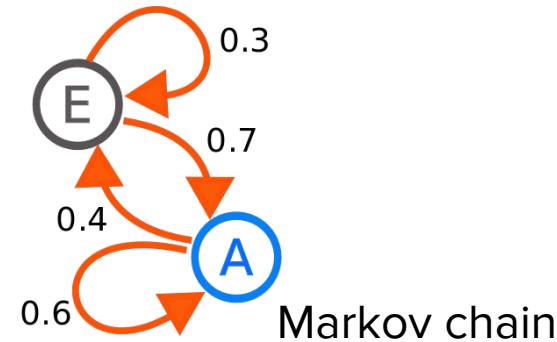
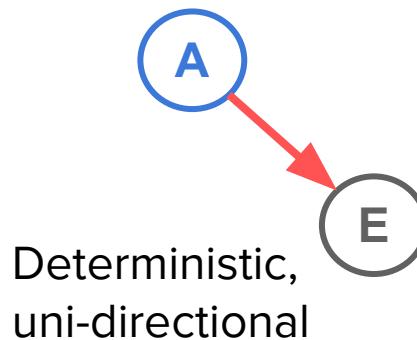


# What if you don't have predefined heuristics?

- How would you solve a problem with
  - large complexity
  - no subject-matter expertise to define any special rules/heuristics that might reduce complexity

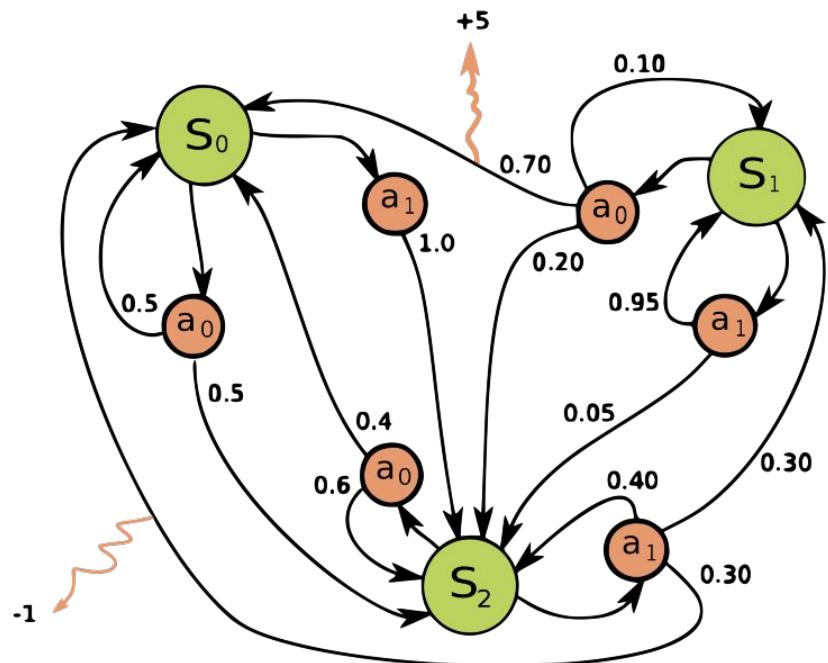
# Let's generalize the problem...

- Rather than imagining state-transitions as a decision tree with deterministic outcomes, let's imagine them as a Markov chain
- In a Markov chain state-transitions are probabilistic



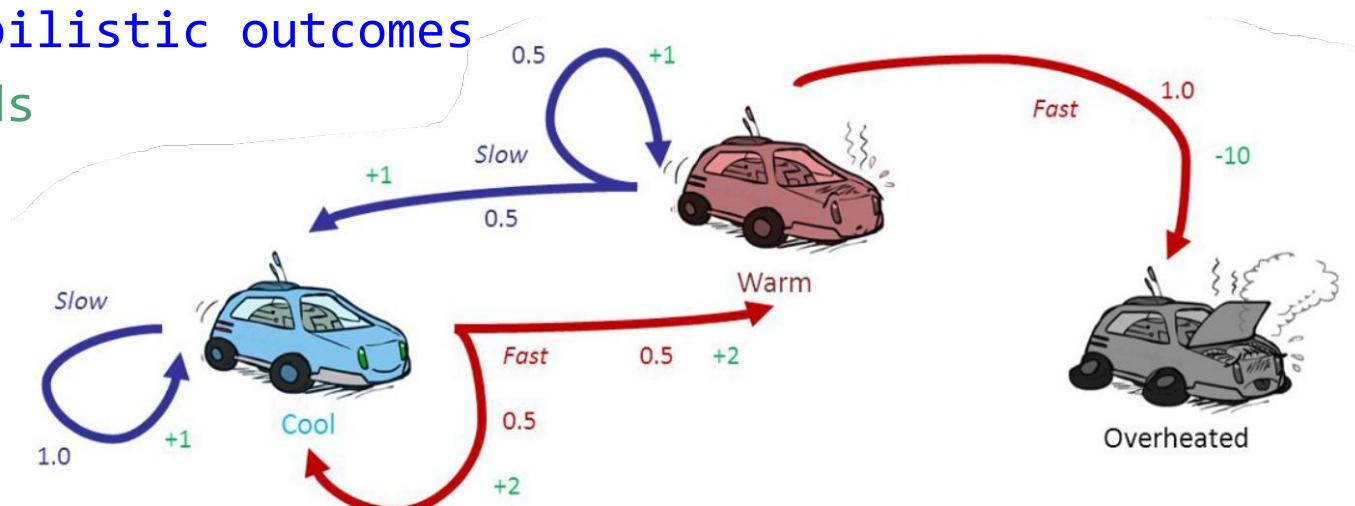
# Markov Decision Process (MDP)

- MDP is an extension of Markov Chains
  - states enable actions
  - actions have
    - probabilistic outcomes
    - rewards



# Markov Decision Process (MDP)

- MDP is an extension of Markov Chains
  - actions have
    - probabilistic outcomes
    - rewards



- What's the best policy for this problem?

## More formally...

- You have a set of states,  $S=\{s_1 \dots s_n\}$
- From each state you have a set of actions,  
 $A(s)=\{a_1 \dots a_n\}$
- Each  $(a|s)$  can lead to some value/reward,  $r$
- You want to find an optimal policy,  $\pi$ , such that
  - you maximize total value from all your actions over time,  $t$ :

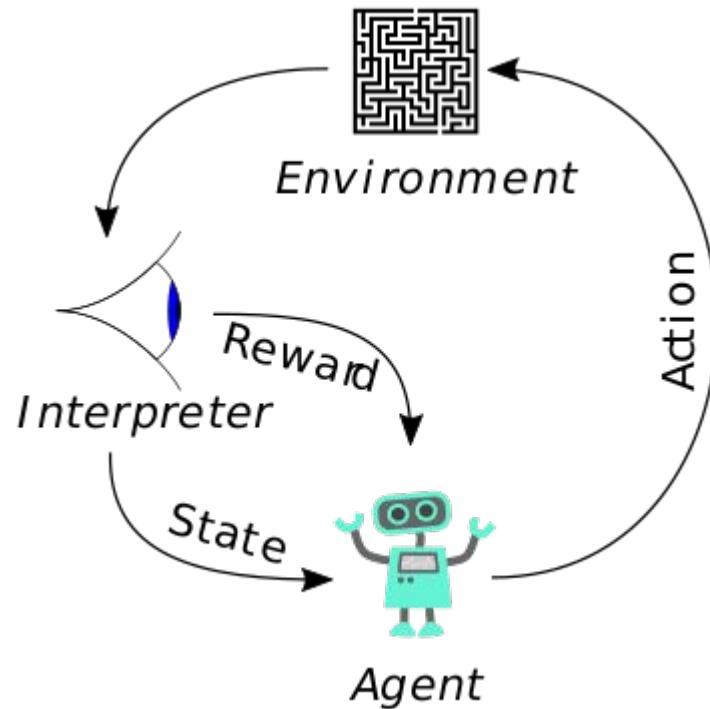
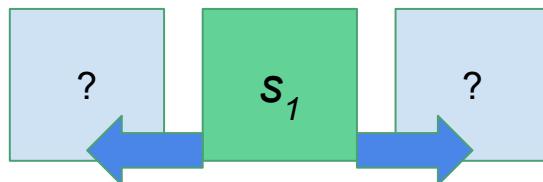
$$\pi \leftarrow \max \left( \sum_{t=0}^{\infty} r_t \right)$$

# When brute-force won't work...

- If your state-space is knowable and small
  - you can do a brute force search to find an optimal policy
- What if your state-space is unknown/unknowable, or just very large?
- What if you don't have predefined heuristics?

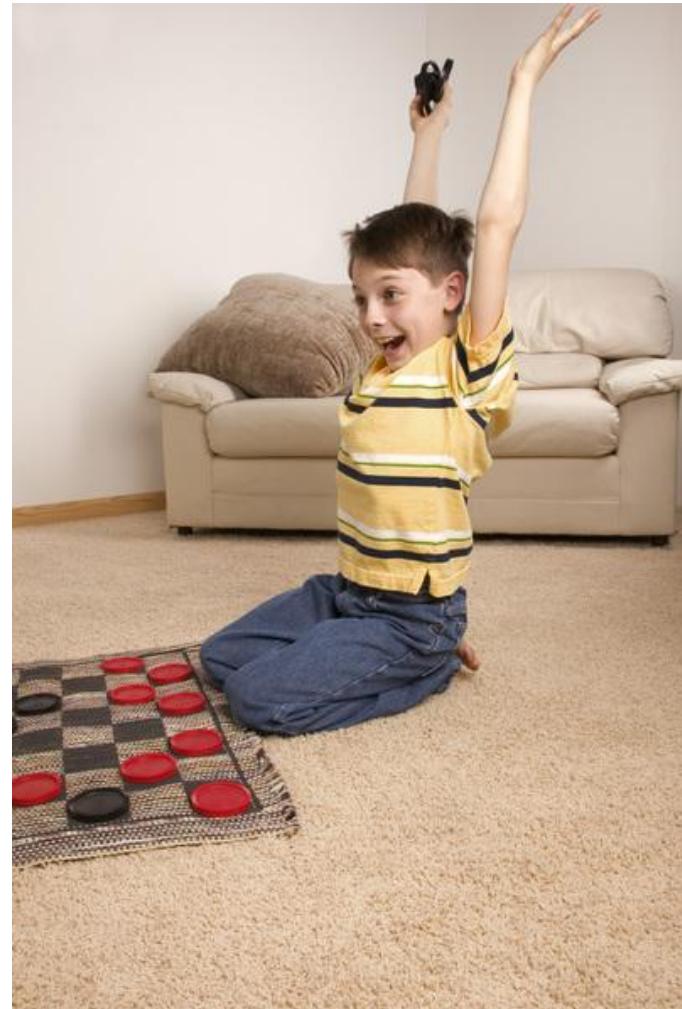
# What would you do?

- You are this Agent
- You are in state  $s_1$
- Do you go left? or right?

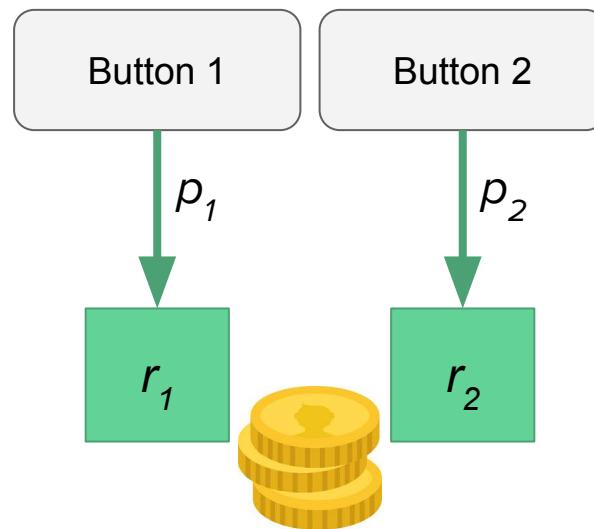


# Trial-and-error learning

- humans do this all the time
  - we try something
    - if it works, we are more likely to do it again
    - if it doesn't, we are less likely to do so



# Let's play an abstract game...



# Let's formalize the problem: Exploration vs Exploitation

- You can do a set of possible actions  $\{a_1, a_2\}$
- Your expected values/rewards are  $Q(a_1)$  and  $Q(a_2)$  for the two actions, respectively
- The **greedy** policy would be to take the action,  $A$ , that has the highest  $Q$ -value
  - $A = \operatorname{argmax} Q(a)$
- But since our  $Q$ -values may not be correct, we want to explore all possible actions and **learn** and update the  $Q$ -values, we should take random actions with some probability,  $\epsilon$

## Let's formalize the problem: Error-driven learning

- After you take action  $A$  you receive a reward,  $r$
- Now we want to learn; i.e., update our expectations for what the value of  $A$  is
  - increase  $Q(A)$  by some proportion,  $\alpha$ , of the difference between actual reward,  $r$ , and the expected reward,  $Q(A)$ 
    - $$Q(A) = Q(A) + \alpha[r - Q(A)]$$
- This is called error-driven learning
- $\alpha$  is the learning rate parameter

# Basic Reinforcement Learning

- Action selection:
  - with a probability of  $\epsilon$ 
    - take a random action
  - with a probability of  $1-\epsilon$ 
    - take action  $A$ , where  $A = \text{argmax } Q(a)$
- Learning:
  - $Q(A) = Q(A) + \alpha[ r - Q(A) ]$

---

$\epsilon$  is the exploration parameter •

$\alpha$  is the learning rate •

# An alternative method for exploration/exploitation

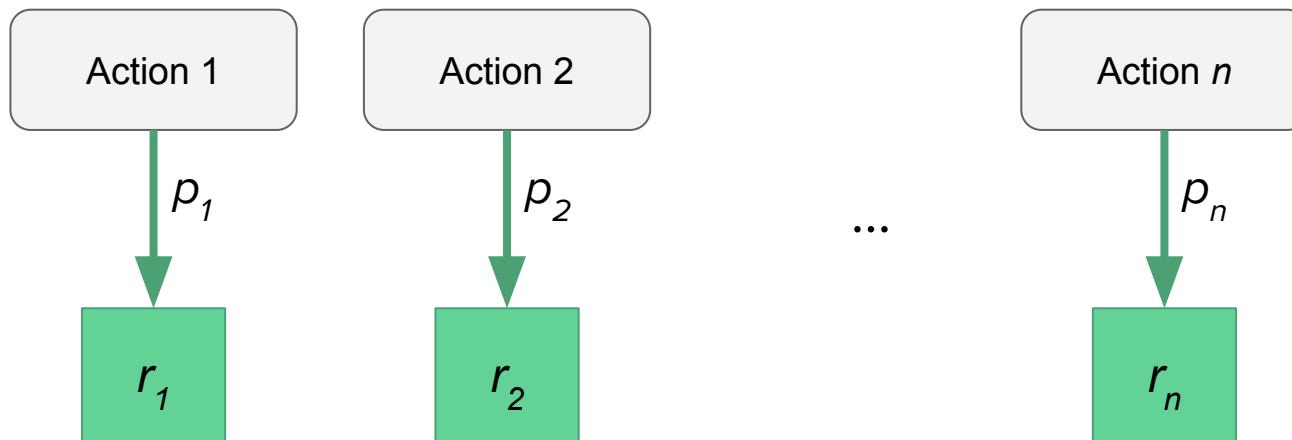
- Action selection:
  - take action  $A$ , where  $A = \operatorname{argmax} [ Q(A) + N(\varepsilon) ]$ 
    - where  $N(\varepsilon)$  is a random value picked from a normal distribution with a standard deviation of  $\varepsilon$
- Learning:
  - $Q(A) = Q(A) + \alpha[ r - Q(A) ]$

---

$\varepsilon$  is the exploration parameter •  
 $\alpha$  is the learning rate •

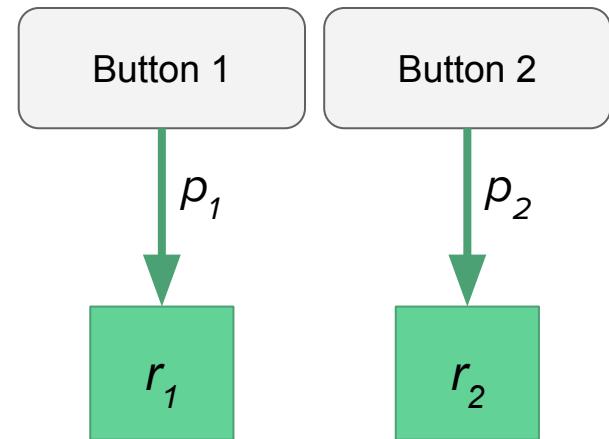
# $n$ -armed bandit

- your agent's job is to maximize its reward over time



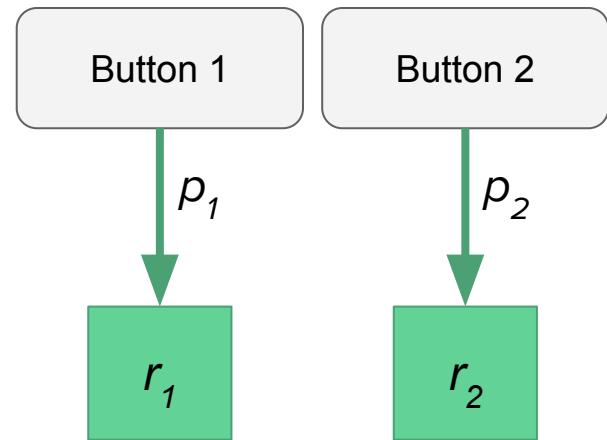
## 2-armed bandit example

- let's say you have 2 buttons
- on step 1, you click Button 1 and get  $r=1$
- assuming learning rate  $\alpha=1.0$ 
  - what would  $Q(\text{Button 1})$  become?
  - what about  $Q(\text{Button 2})$ ?
- what would a greedy ( $\epsilon=0$ ) agent do on step 2?
- assume you get  $r=0$  on step 2
  - now what?



## 2-armed bandit example

- let's say you have 2 buttons
- on step 1, you click Button 1 and get  $r=1$
- assuming learning rate  $\alpha=0.1$ 
  - what would  $Q(\text{Button 1})$  become?
  - what about  $Q(\text{Button 2})$ ?
- what would  $\epsilon=0.5$  agent do on step 2?
- if Button 1 is clicked again, and  $r=0$ , what are the new  $Q$ -values?



# Reinforcement Learning vs Minimax

- Minimax seems optimal
  - it uses brute force search to find *THE* best move every time
  - it will never lose
  - but...
- In the real world
  - a reinforcement learner playing tic-tac-toe will have way more wins than minimax
  - why do you think this might be?

# Assignment 3

- add reinforcement learning agent code to the supplied jupyter notebook and run the notebook
- *details on blackboard*

# Reinforcement Learning (part 2)

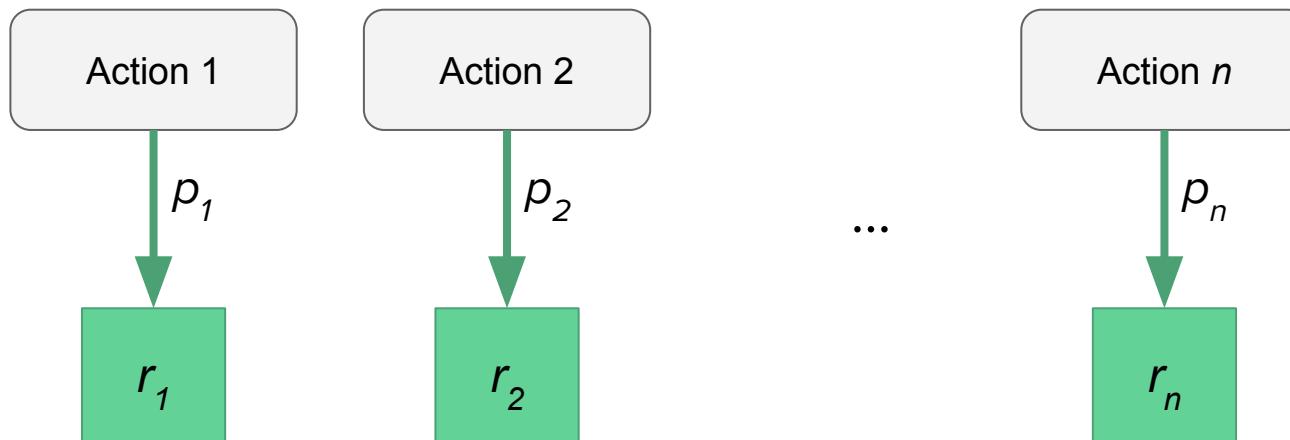
---

CS420 - Lecture 4  
Instructor: Dr. V



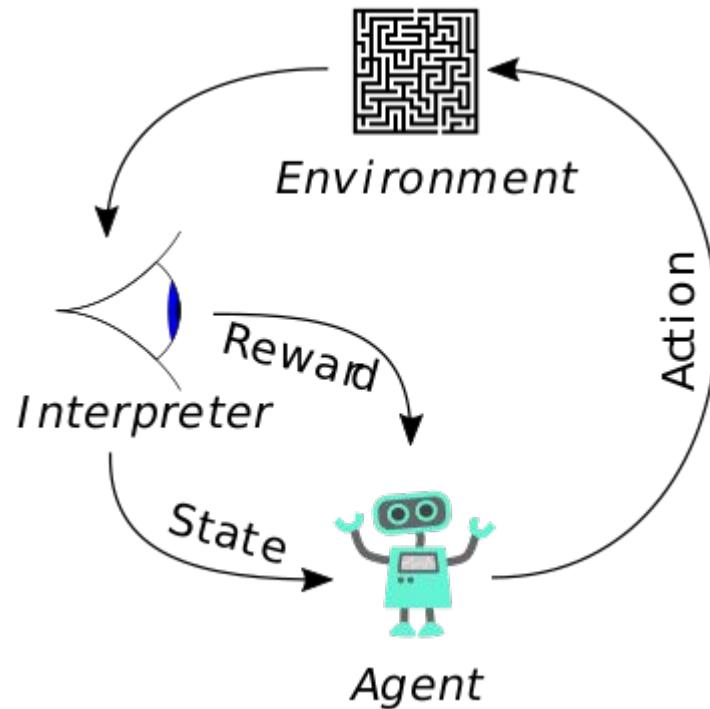
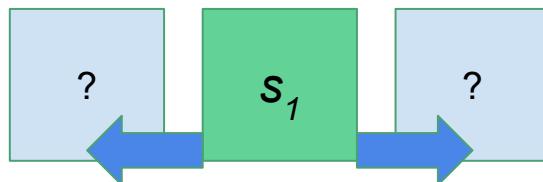
# $n$ -armed bandit

- your agent's job is to maximize its reward over time



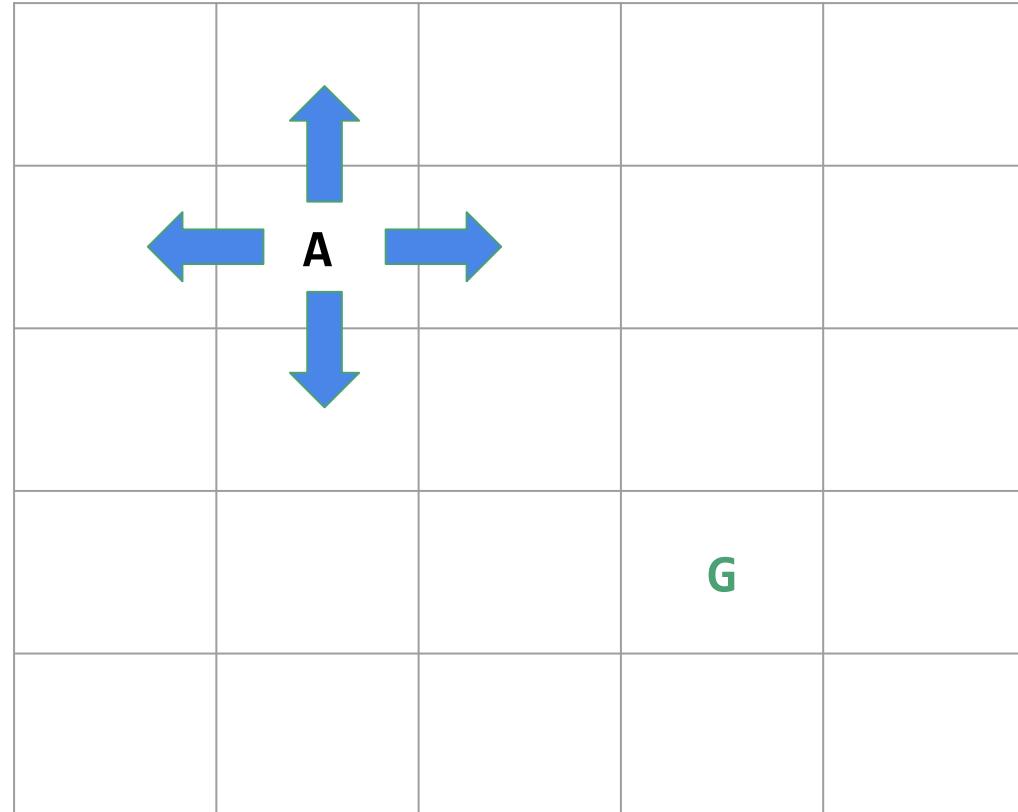
# What would you do?

- You are this Agent
- You are in state  $s_1$
- Do you go left? or right?



# Imagine a navigation environment

- Agent, A, gets a reward if it reaches its goal, G
- 4 possible actions from each square
  - it's as if each square is a separate  $n$ -armed bandit problem



# So how do we design agent memory?

- For a single-state agent ( $n$ -armed bandit)
  - we had an **array** of length  $n$  to store our  $Q$ -values
- So how would we store  $Q$ -values for an agent that can take  $n$  possible actions from each of  $m$  possible states?

- in this example
  - 25 states
  - 4 actions

		A	
		G	

Q-table initialised at zero

	UP	DOWN	LEFT	RIGHT
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0
7	0	0	0	0
8	0	0	0	0

After few episodes

	UP	DOWN	LEFT	RIGHT
0	0	0	0	0
1	0	0	0	0
2	0	2.25	2.25	0
3	0	0	5	0
4	0	0	0	0
5	0	0	0	0
6	0	5	0	0
7	0	0	2.25	0
8	0	0	0	0

Eventually

	UP	DOWN	LEFT	RIGHT
0	0	0	0.45	0
1	0	1.01	0	0
2	0	2.25	2.25	0
3	0	0	5	0
4	0	0	0	0
5	0	0	0	0
6	0	5	0	0
7	0	0	2.25	0
8	0	0	0	0

# Storing Q-values

- For every state, for every action in that state, you need to store a  $Q$ -value (a floating point value)
  - you can use a matrix (i.e., 2-dimensional array)
  - or a dictionary (i.e., hashmap)

```
{
```

```
    state1: [0.0, 0.0, 0.2, 0.0],
```

```
    state2: [0.0, 1.0, 0.0, 0.0],
```

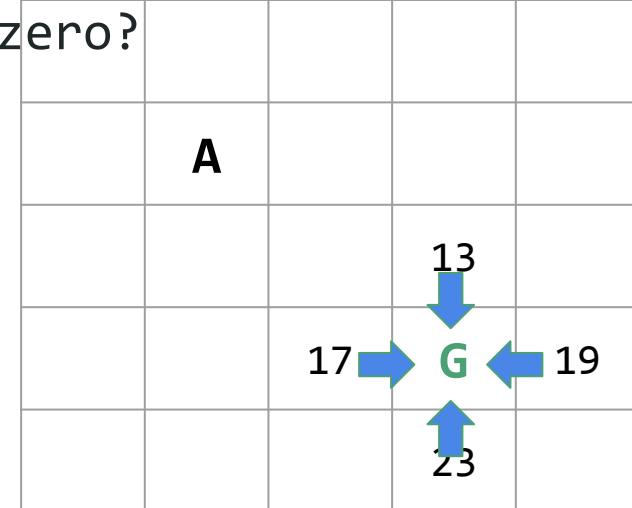
```
    state3: [0.0, 0.0, 0.0, 0.5],
```

```
    ...
```

```
}
```

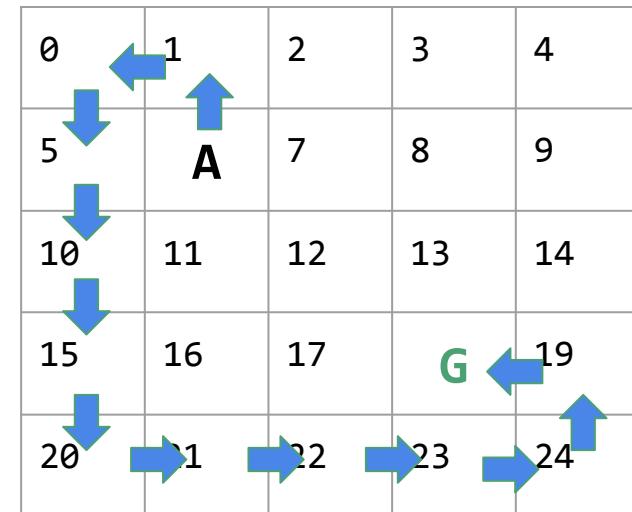
# How do we assign credit for reaching reward?

- If we still do what we did for the  $n$ -armed bandit RL
  - state-action pairs 13-down, 17-right, 19-left, and 23-up will get rewarded
    - but what about all other state-action pairs (SAs)?
    - do all other  $Q$ -values stay at zero?



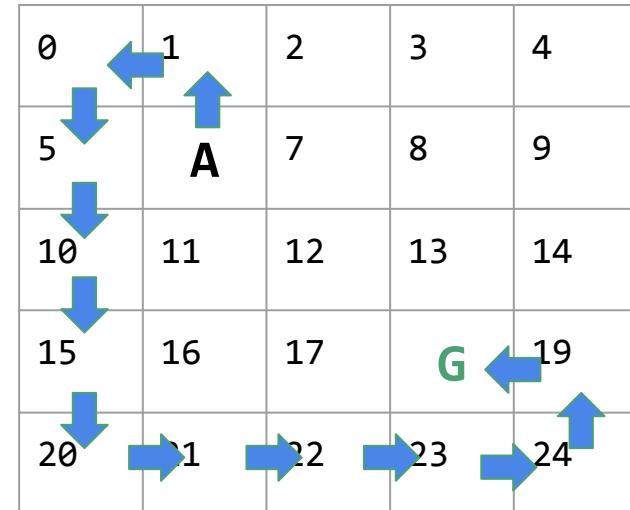
# Credit attribution problem

- Imagine agent is taking random actions
  - happens to take the blue actions shows
  - reaches space 19, happens to take a left
  - reaches the goal, gets a reward
- What should agent learn?
  - that 19-left is good?
  - that all SAs it took are good?



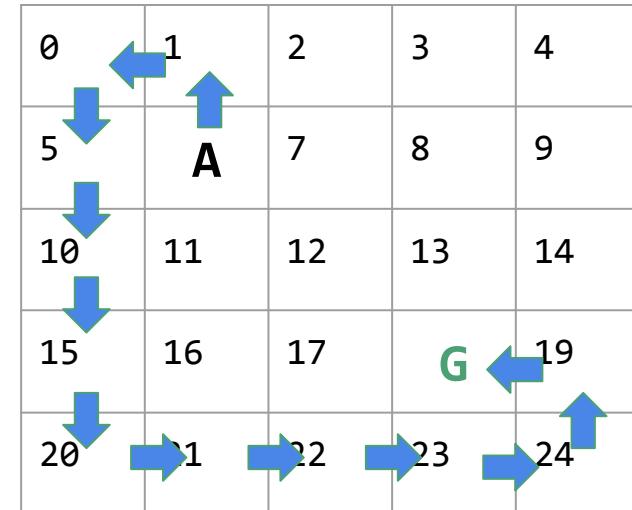
# Credit attribution problem

- What should agent learn about the SAs it has taken?
  - is 19-left is good?
  - is 24-up good?
  - is 6-up good?
  - is 1-left good?



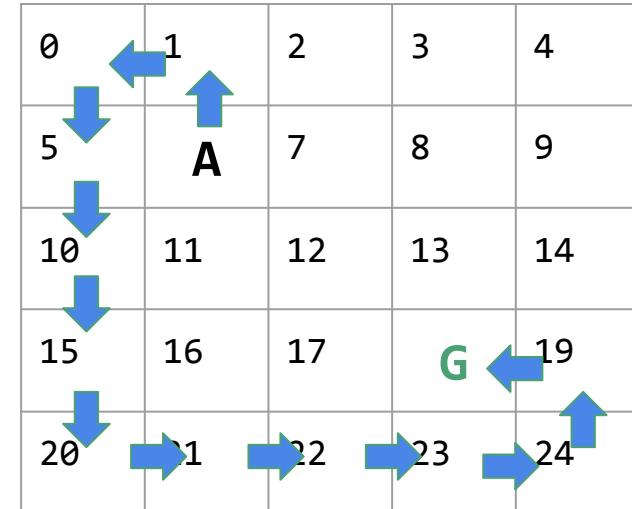
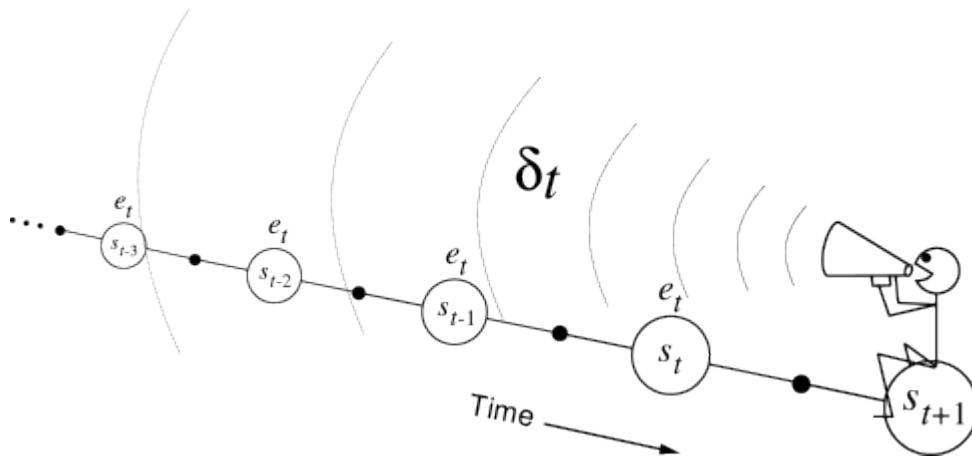
# Credit attribution problem

- There are two main ways to propagate reward back to all actions that lead to it
  - looking into the past
    - Eligibility traces
  - looking into the future
    - Temporal-Difference learning
      - Q-learning



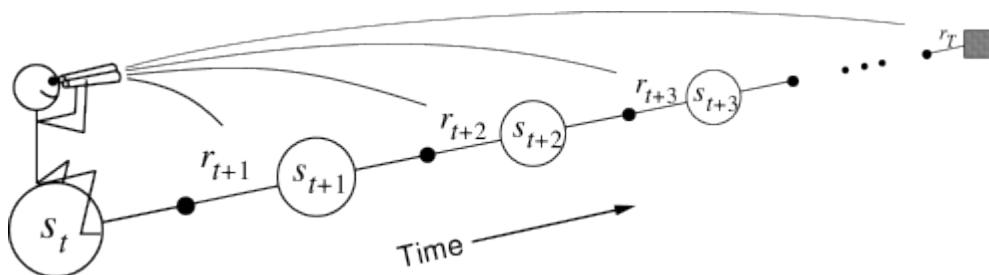
# RL with Eligibility Traces

- Keep a FIFO queue of your SAs  
[6↑, 1←, 0↓, 5↓, ..., 23→, 24↑, 19←]
- When you get a reward, propagate it back to all SAs in the queue, giving lower rewards to earlier SAs



# Q-Learning

- When you get a reward, assign it only to  $(S_x, A)$  that took you there
  - When, later, another  $(S_y, A)$  leads to  $S_x$ , its Q-value should be increased by a proportion of maximum Q-value from  $S_x$



0	1	2	3	4
5	A	7	8	9
10	11	12	13	14
15	16	17	G	19
20	21	22	23	24

# Basic error-driven vs Q -learning

- Error-driven learning:

- $$Q(s_t, a_t) += \alpha [ r_{t+1} - Q(s_t, a_t) ]$$

- Q-learning:

- $$Q(s_t, a_t) += \alpha [ r_{t+1} + \gamma \cdot \max Q(s_{t+1}, a) - Q(s_t, a_t) ]$$

---

$\gamma$  is the future-value discount factor •  
 $\alpha$  is the learning rate •

# Lab 4

- add standard, q-learning, and eligibility traces RL agents to the supplied jupyter notebook and run the notebook
- details on blackboard

# Assignment 4 (due before next lecture)

- Each team
  - is responsible for adding the following to the exam
    - 2 questions based on Lecture 1
    - 3 questions based on Lecture 2
    - 5 questions based on Lectures 3 and 4
  - should designate an ambassador, who will
    - work with other teams to make sure your team's questions are unique
    - present 1 veto per section during next lecture
- Questions must be multiple choice, with answer key, and "Shuffle option order"
- Add your questions directly to shared form, AND paste your team's questions on blackboard

# Supervised Learning

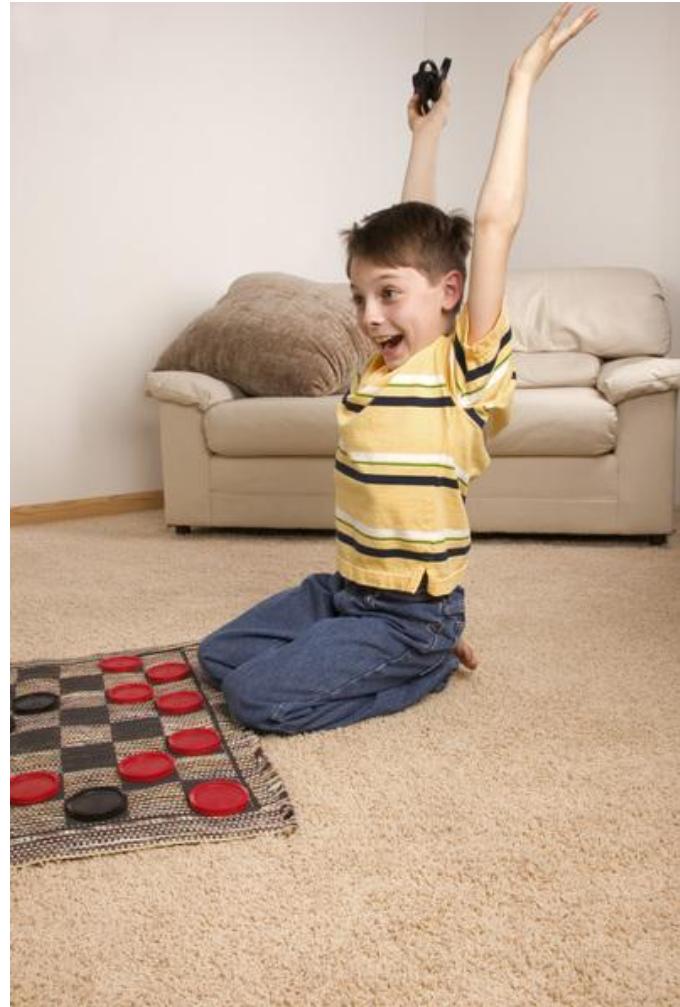
---

CS420 - Lecture 4  
Instructor: Dr. V



# Reinforcement Learning

- Relevant for:
  - Trial-and-error problems
- How it works:
  - try something
  - observe feedback
  - if the reward is high, do more of it
  - if reward the is low, do less of it



# Supervised Learning

- Relevant for:
  - **Labeling data**
- How it works:
  - observe input features
  - guess label
  - observe correct label
  - next time there are similar input features,  
guess the correct label



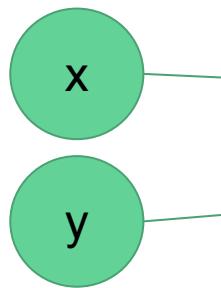


02

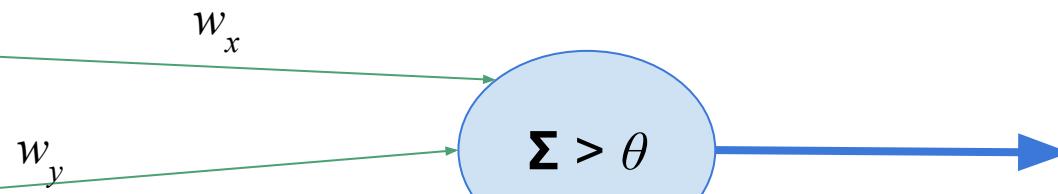
# SUPERVISED LEARNING

# Perceptron

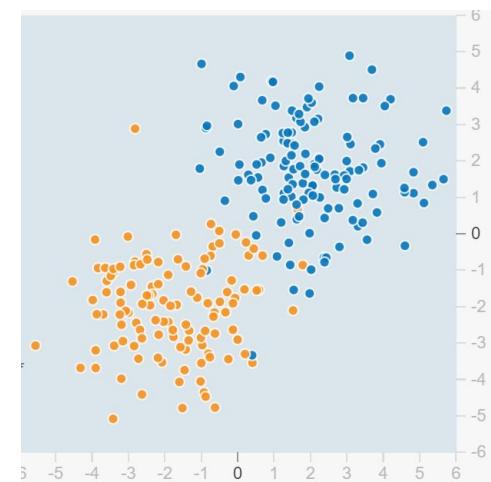
Inputs



Perceptron

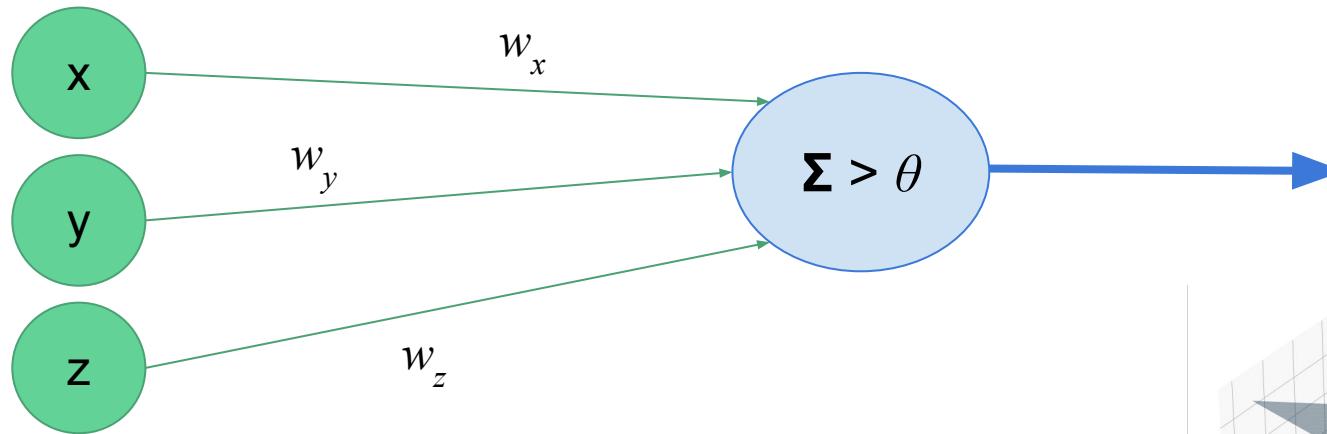


Output



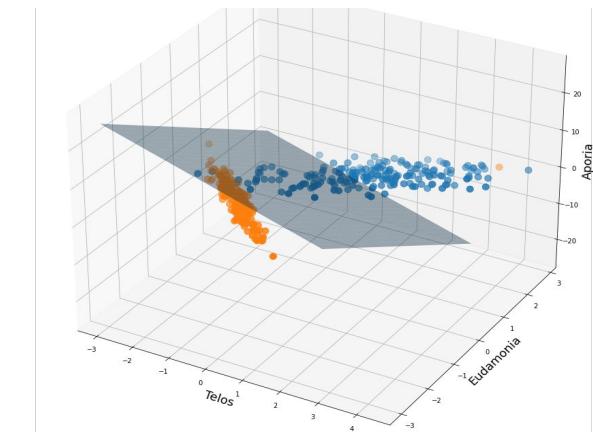
# Perceptron

Inputs



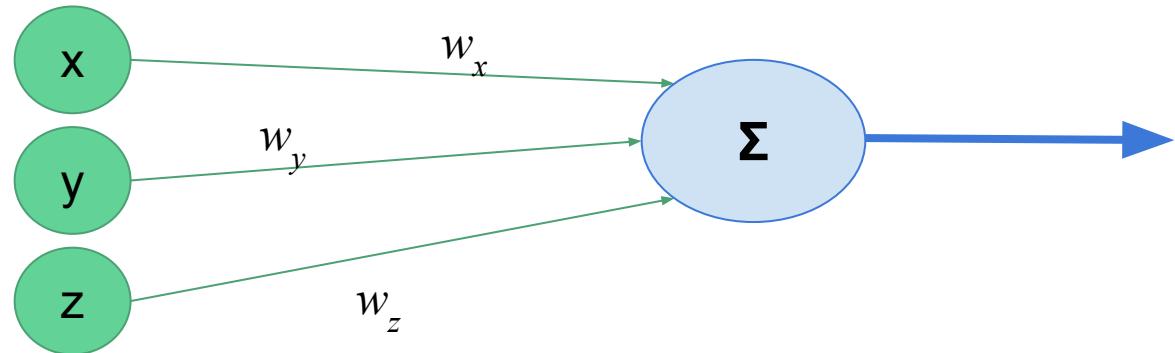
Perceptron

Output



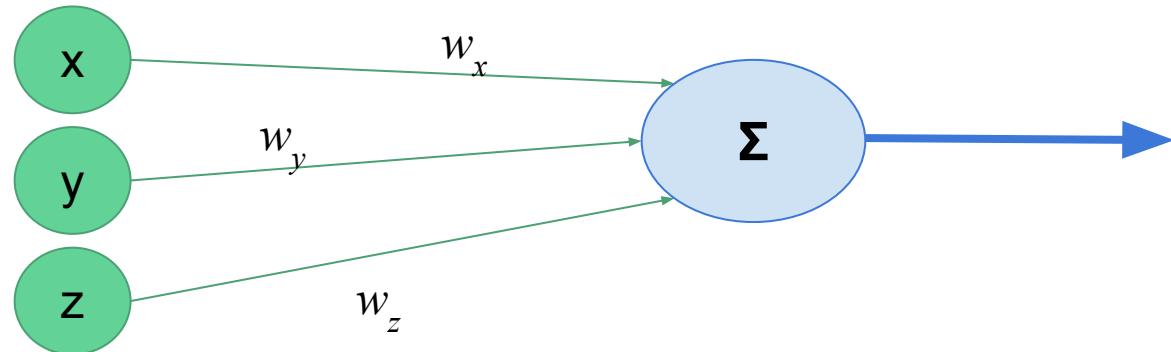
# How does a perceptron learn?

- if the  $\Sigma$  output is, let's say 0.3, but supervisor says it should have been 1.0, the error is +0.7
  - so all the weights need to get bumped up by some proportion of 0.7
  - the weights that brought the most signal get bumped up the most



# How does a perceptron learn?

- if the  $\Sigma$  output is, let's say 0.3, but supervisor says it should have been 0.0, the error is -0.3
  - so all the weights need to get pushed down by some proportion of 0.3
  - the weights that brought the most signal get pushed down the most



## More formally: Delta learning rule

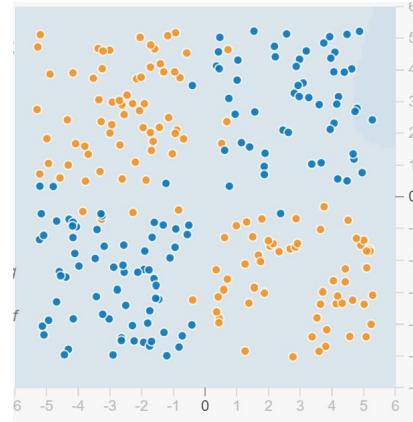
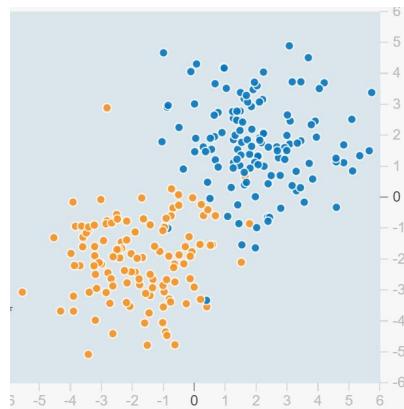
- After you make a prediction
- Update the weight,  $w$ , from each input node,  $i$ 
  - increase  $w_i$  by some proportion,  $\eta \cdot x_i$ , of the difference between the target output,  $t$ , and the predicted output,  $o$ 
    - $w_i = w_i + \eta[ t - o ]x_i$
- $\eta$  is the learning rate parameter
- $x_i$  is the input value for node  $i$

## How does a perceptron learn?

- <https://www.youtube.com/watch?v=khUVIZ3MON8>
- <https://www.youtube.com/watch?v=Ilg3gGewQ5U>

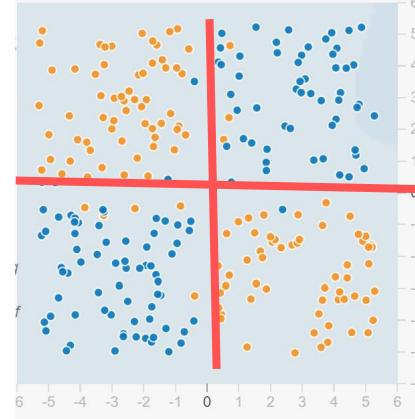
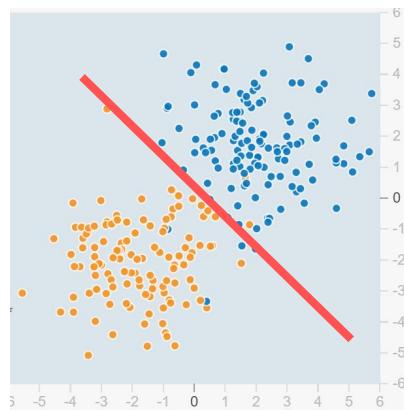
# Limitations of a perceptron

- Data must be linearly separable
  - A single perceptron can draw a single boundary in your feature-space
  - For lots of lines (for non-linearly separable data) you need more neurons connected in layers



# Limitations of a perceptron

- Data must be linearly separable
  - A single perceptron can draw a single boundary in your feature-space
  - For lots of lines (for non-linearly separable data) you need more neurons connected in layers



# Lab 5

- use a perceptron to
  - a. classify which flower instances are Setosa Irises (as opposed to Versicolor or Virginica)
  - b. classify handwritten digits
- follow in-class instructions
- submit assignment notebooks on blackboard

# Assignment 5 (due before next lecture)

- Each team
  - is responsible for adding 4 questions based on Lectures 5 to the exam
- Team ambassador will
  - work with other teams to make sure your team's questions are unique
- Questions must be
  - multiple choice
  - with answer key
  - "Shuffle option order" (if it makes sense)
- Add your questions directly to shared form, AND paste your team's questions on blackboard

# Neural Networks

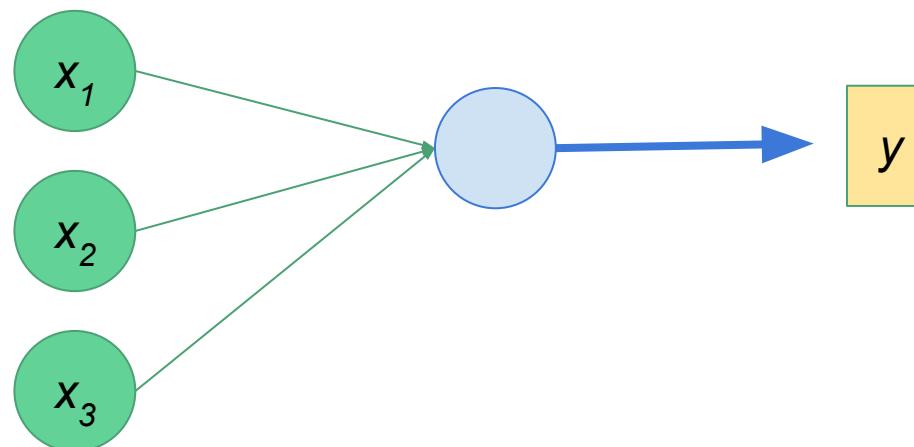
---

CS420 - Lecture 4  
Instructor: Dr. V



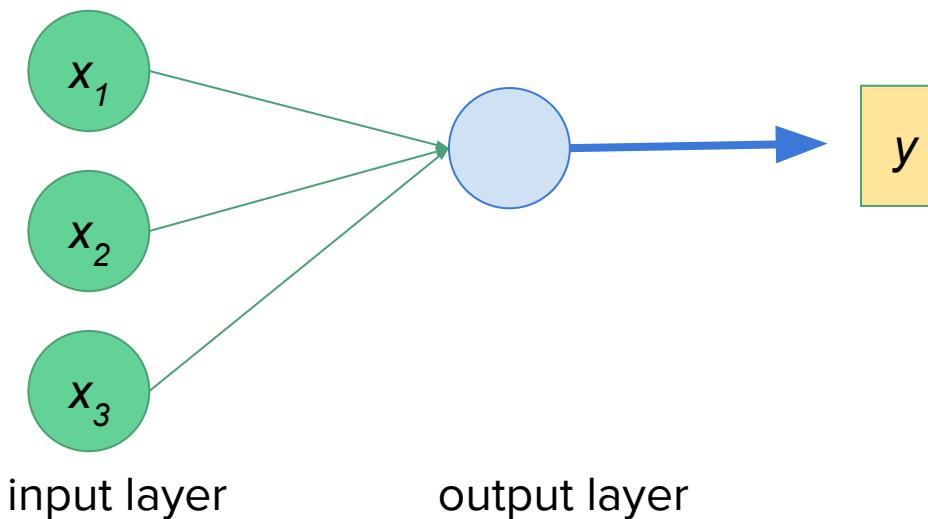
# Perceptron

- Each perceptron is a single neuron
  - it has some numeric inputs,  $x$ , and it produces a numeric output,  $y$



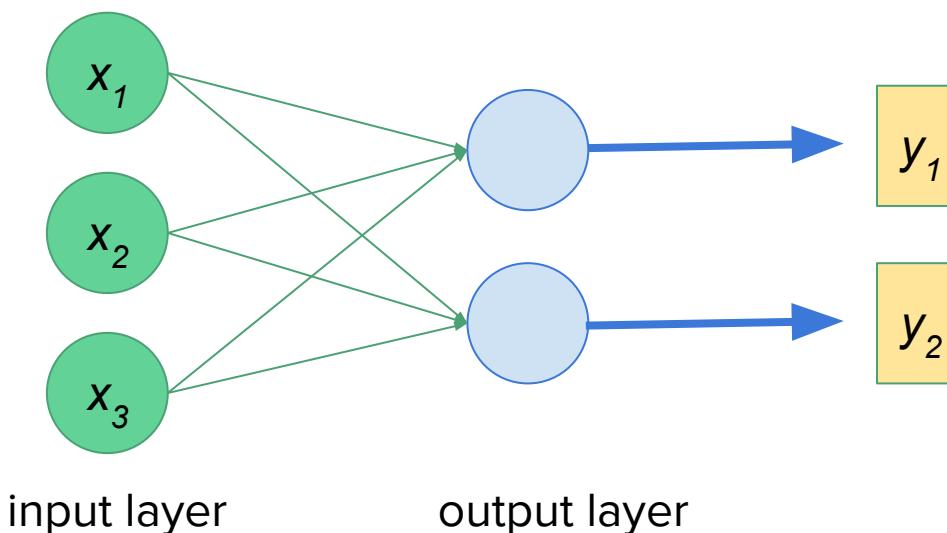
# Neural Networks

- We can technically call a perceptron a neural network, rather than a neuron, since the input nodes can be thought of as neurons, as well



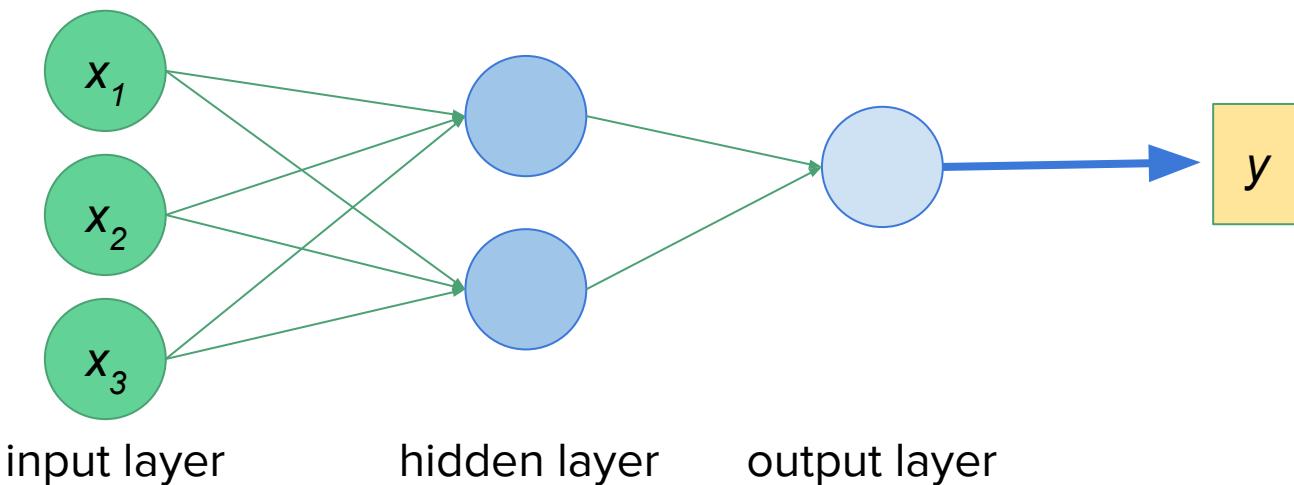
# Neural Networks

- If you needed more than a single output, you could add more perceptrons, all connected to and activated by the same input nodes



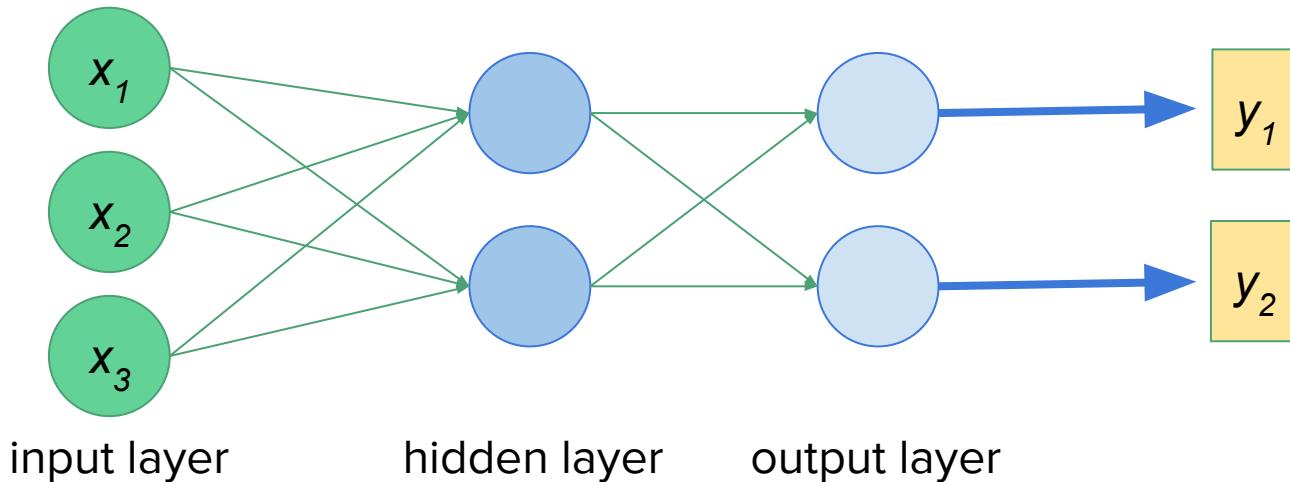
# Feed-forward Neural Networks

- A **feed-forward** neural network is made up of multiple perceptrons in successive "**layers**", where outputs from prior layer neurons is used as inputs for neurons in next layer



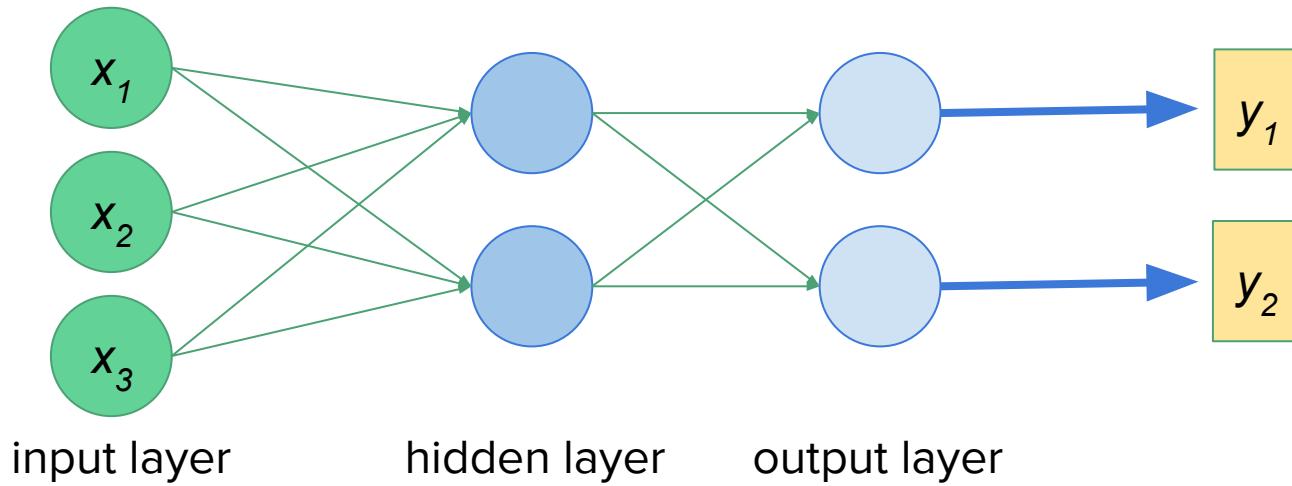
# Feed-forward Neural Networks

- In a feed-forward neural net, activation is spread forward from input nodes to each successive layer, until output is produced



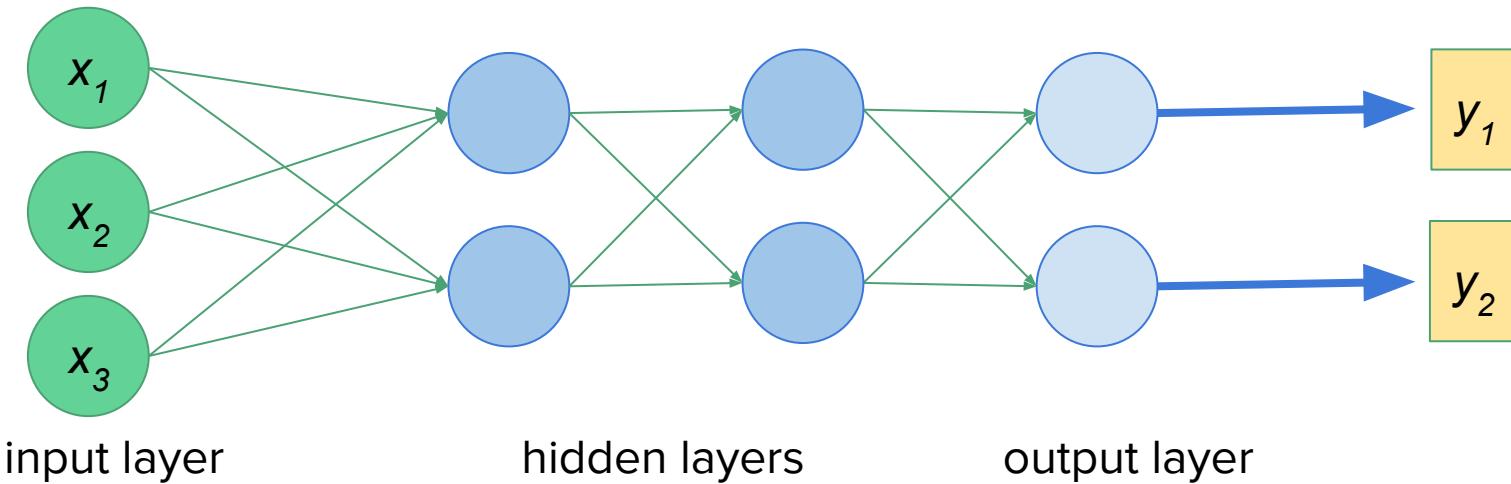
# Feed-forward Neural Networks

- Most common types of Artificial Neural Nets (ANNs) employed for supervised learning are **feed-forward** nets with **fully interconnected** successive layers



# Deep Neural Networks

- Deep nets are feed-forward networks with more than one hidden layer



# Terminology

- Machine Learning (ML)
  - a subfield of AI concerned with trainable models/agents
- Deep Learning (DL)
  - a subfield of ML concerned with using deep neural networks

# Shape of the network

- Why would you want more layers?
- Why would you want more neurons in a layer?
  
- Why would you want less layers?
- Why would you want less neurons in a layer?

<https://playground.tensorflow.org/>



# NEURAL NETWORKS



# How does a neural network learn?

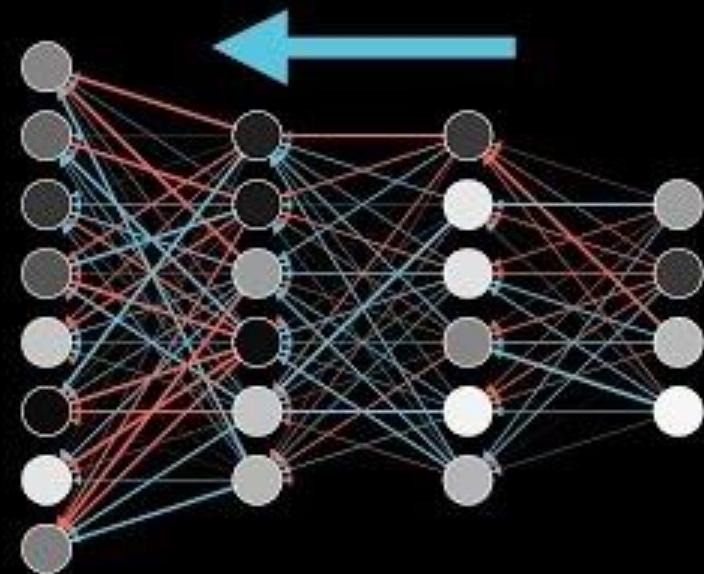
- We know how every neuron in output layer learns:
  - Delta rule
    - the error between perceptron output and target output is used to adjust input weights
- How do the neurons in earlier layers learn?
  - Backpropagation
    - the error each successive layer is passed back to the prior layer



04

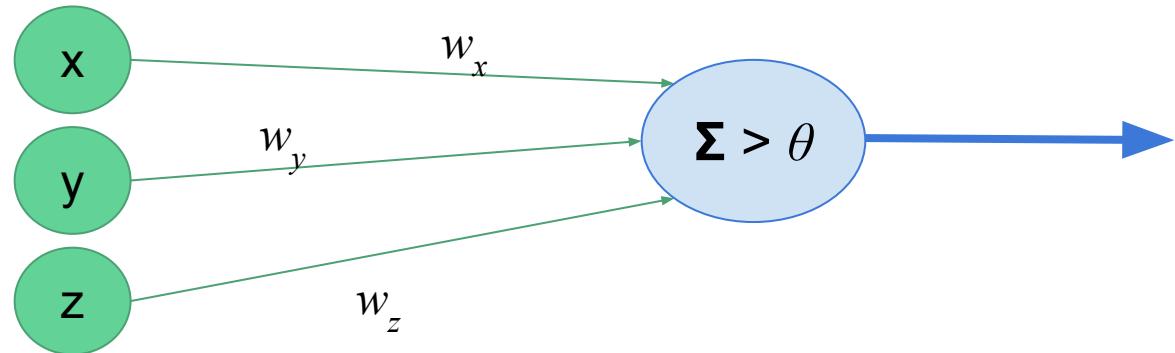
# BACKPROPAGATION AND OPTIMIZATION

# Backpropagation

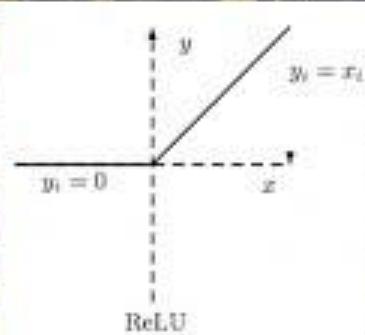


# Activation functions

- After we gather the sum of input activations, we output  $y$  based on some activation function; e.g.,
  - step function
  - linear
  - **ReLU**
  - leaky ReLU
  - sigmoid
  - tanh



# WHICH ACTIVATION FUNCTION SHOULD I USE?



# Recommendations

- for input where neighboring features are related
  - use convolutional layers
- for activation of hidden layers
  - use relu
- for activation of output layer
  - for classification, use softmax
  - for regression, use a linear function
- to calculate error
  - for classification, use categorical cross-entropy
  - for regression, use mean-abs (or mean-sqr) error

# Popular python neural network libraries

- Tensorflow / keras
- PyTorch
- NeuroLab
- ffnet
- Scikit
- Lasagne
- pyrenn
- MXNet

# Assignment 6 (due before next lecture)

- Each team
  - is responsible for adding 4 questions based on Lectures 6 to the exam
- Team ambassador will
  - work with other teams to make sure your team's questions are unique
- Questions must be
  - multiple choice
  - with answer key
  - "Shuffle option order" (if it makes sense)
- Add your questions directly to shared form, AND paste your team's questions on blackboard

# Lab 6

- use tensorflow to classify
  - MNIST handwritten digits
  - compare against perceptron accuracy
- details on blackboard

# tensorflow.keras

```
from tensorflow import keras
model = keras.models.Sequential([
    keras.layers.Dense(128, activation='relu'), # hidden layer 1
    # add more hidden layers here, if you like
    keras.layers.Dense(10, activation='softmax') # output layer
])

model.compile(
    # learn via backpropagation (adam is a great algorithm for bp)
    optimizer='adam',
    # calculate error based on categorical cross-entropy
    loss=keras.losses.sparse_categorical_crossentropy,
    # eval model based on accuracy
    metrics=['accuracy']
)
```

# Unsupervised Learning

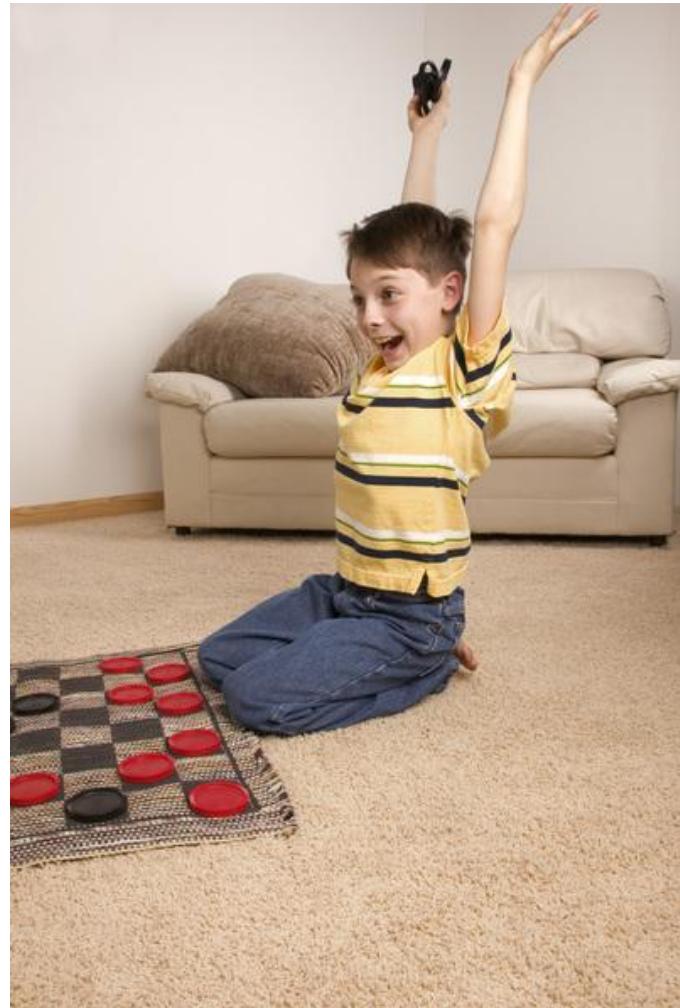
---

CS420 - Lecture 7  
Instructor: Dr. V



# Reinforcement Learning

- Relevant for:
  - Trial-and-error problems
- How it works:
  - try something
  - observe feedback
  - if the reward is high, do more of it
  - if reward the is low, do less of it



# Supervised Learning

- Relevant for:
  - **Labeling data**
- How it works:
  - observe input features
  - guess label
  - observe correct label
  - next time there are similar input features,  
guess the correct label

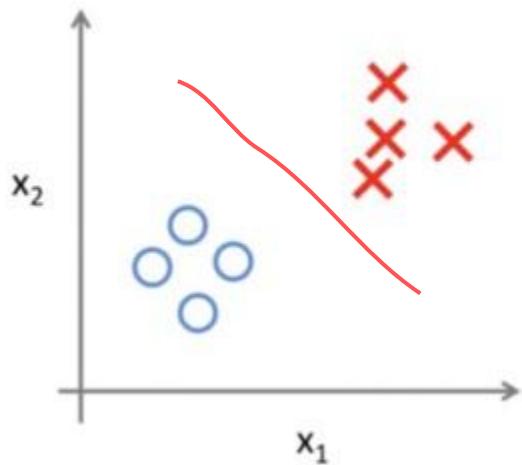


# Unsupervised Learning

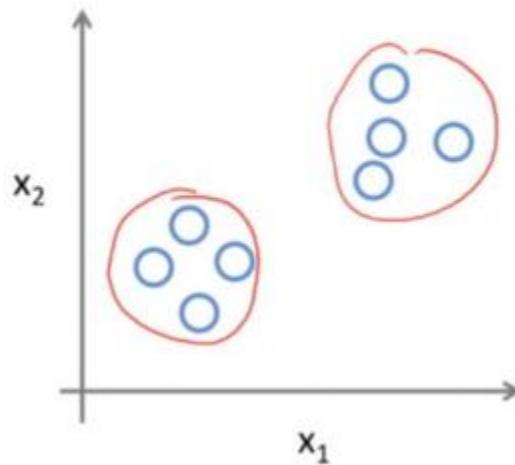
- Relevant for:
  - **Everything else**
    - No corrections
    - No reward feedback
- How it works:
  - observe inputs
  - learn something about them; e.g.,
    - similarity to other inputs, prototype



## Supervised Learning



## Unsupervised Learning



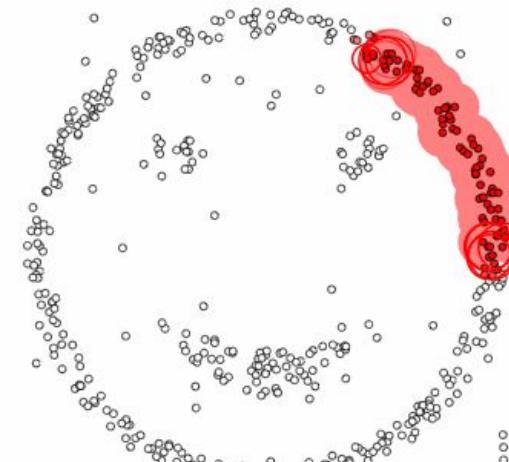
# Examples of unsupervised learning

- Dimensionality reduction
  - reducing the number of input features in a dataset (e.g., PCA, SVD, Autoencoders)
- Anomaly detection
  - detecting instances that are very different from the norm (e.g., Isolation Forest, Minimum Covariance Determinant)
- Clustering
  - grouping similar instances into clusters (e.g., k-means, hierarchical clustering)

# Clustering

- There are many clustering algorithms
  - k-means
  - mean-shift clustering
  - DBSCAN
  - BIRCH
  - hierarchical clustering
  - gaussian mixture models

epsilon = 1.00  
minPoints = 4



Restart



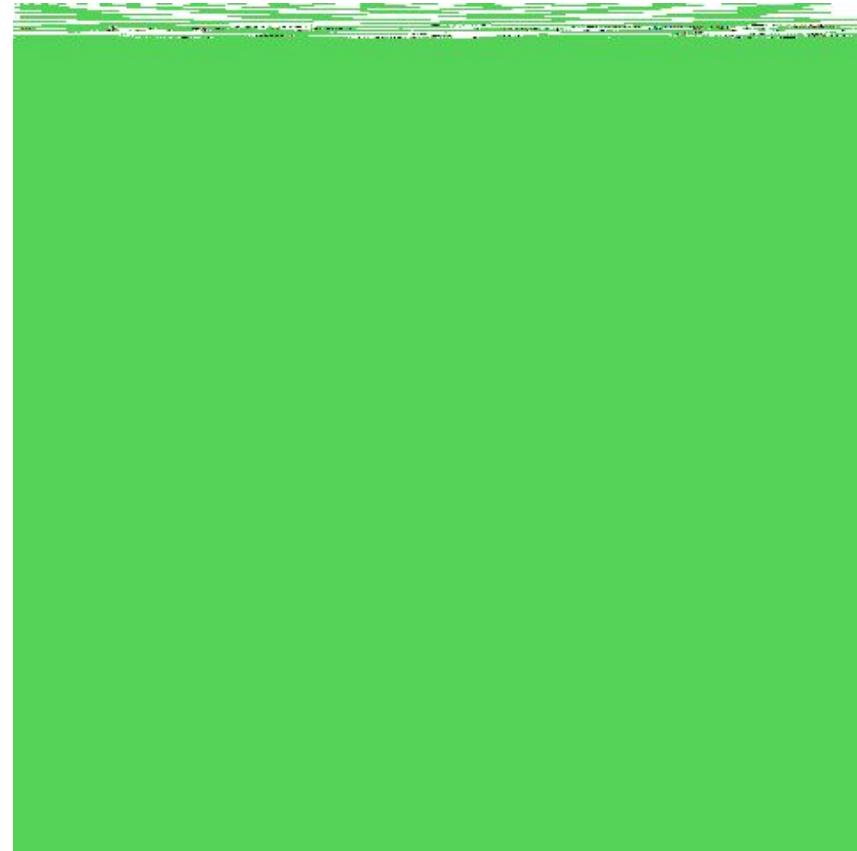
Pause

# K-means Clustering

```
## K-means Clustering Pseudocode
1. Choose the number of clusters, K
2. Place the centroids ( $c_0 \dots c_{k-1}$ ) randomly
3. Repeat until centroids stop changing:
    # assign data points to clusters
    for each data point:
        - find the nearest centroid ( $c_0 \dots c_{k-1}$ )
        - assign the point to that cluster (0 ... k-1)
    # change centroids to reflect new clusters
    for each cluster j in (0 ... k-1):
        -  $c_j$  = mean of all points assigned to that cluster
```

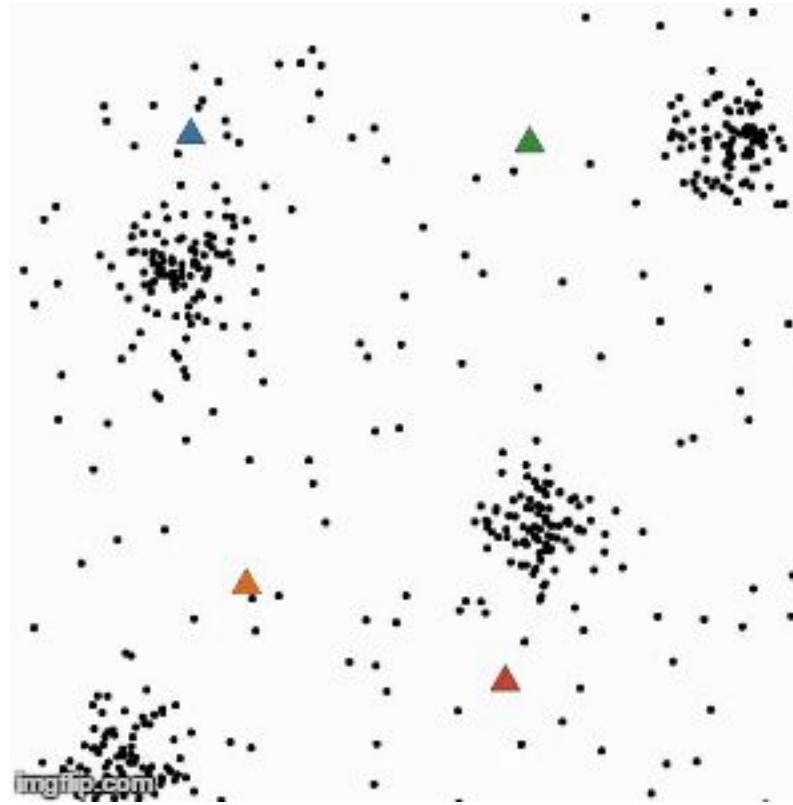
# K-means Clustering

k = 3



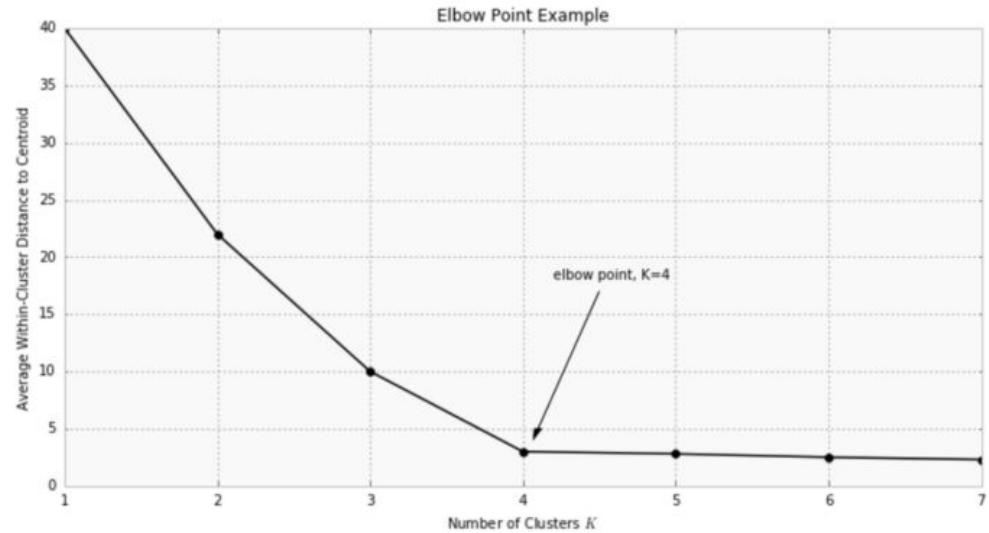
# K-means Clustering

$k = 4$



# How to choose $K$ ?

- most popular: elbow method
  - once average within-cluster distance stops significantly changing, stop increasing  $K$





06

# UNSUPERVISED LEARNING

# Assignment 7 (due before next lecture)

- Each team
  - is responsible for adding 2 questions based on Lectures 7 to the exam
- Team ambassador will
  - work with other teams to make sure your team's questions are unique
- Questions must be
  - multiple choice
  - with answer key
  - "Shuffle option order" (if it makes sense)
- Add your questions directly to shared form, AND paste your team's questions on blackboard

# Lab 7

- create k-means learning algorithm
  - cluster irises using k-means
- submit notebook on blackboard

# NLP

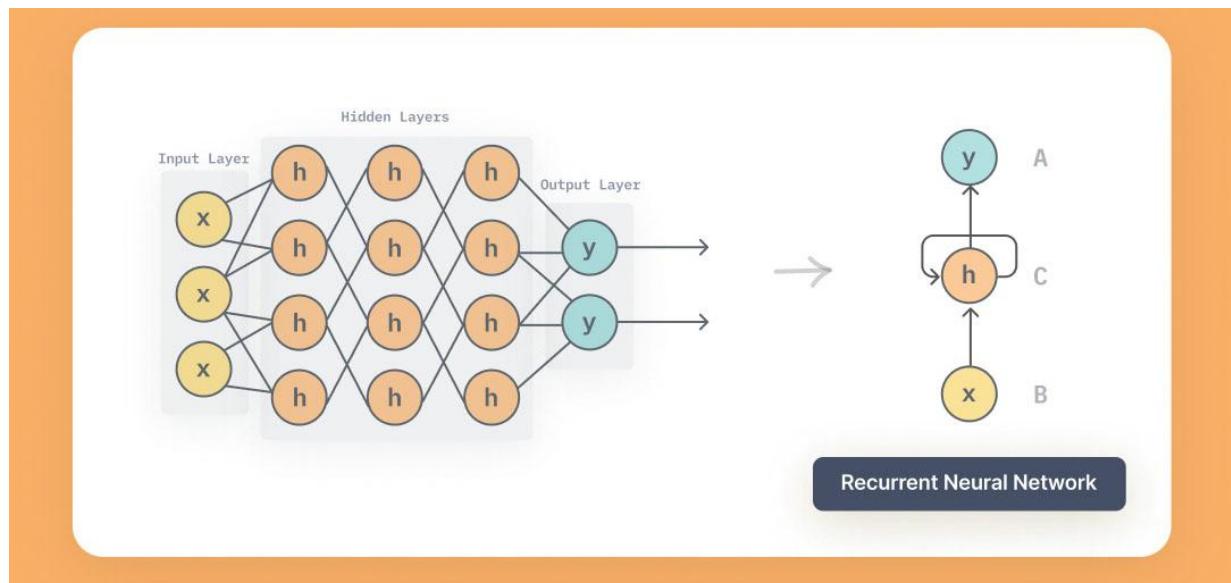
---

CS420 - Lecture 8  
Instructor: Dr. V



# RNN

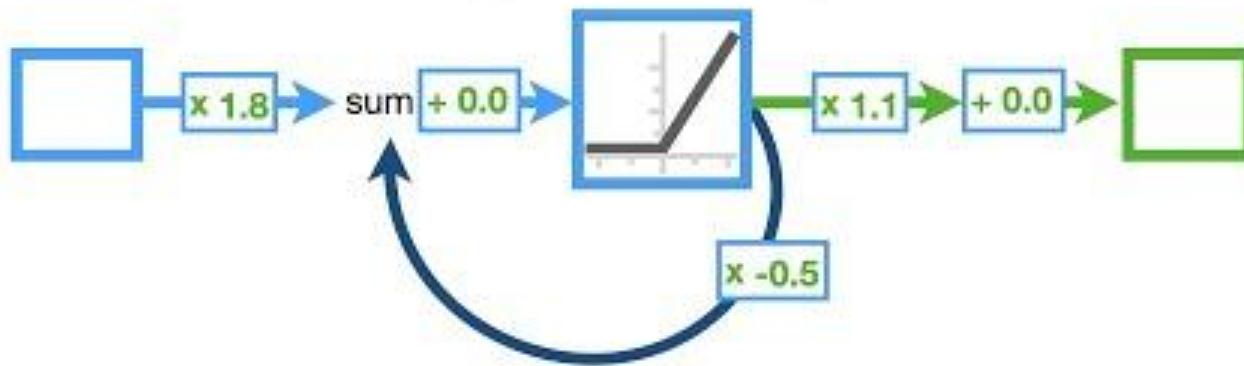
- in a recurrent neural net, the output from the hidden layer is used as the input back to the hidden layer



# RNN

- RNN's can learn sequences
  - assume you want to learn which words follow which other words, so that you can predict text
    - input: "hello my name is Inigo Montoya"
  - "hello" is the first input, target output is "my"
  - RNN learns that "hello" output should be "my"
  - now "my" becomes the input, target output is "name", but the hidden layer output from "hello" is ALSO part of the input
  - etc.

# Recurrent Neural Networks (RNNs)...



...Clearly Explained!!!

# RNN

- is RNN supervised or unsupervised learning?

# LSTM (long-term short-term memory) neural nets

- LSTM neural net is like RNN, but older memories have less influence as inputs into the hidden layers than newer memories

# NLP

- there are many aspects to natural language processing in AI
  - predicting which words/terms are next in a sequence
  - grammar parsing
  - finding common terms
  - figuring out how what words/terms are most important/representative of a document
  - figuring out how similar or related two words are
  - figuring out how similar or related two sentences or documents are

## tf-idf

- which words/terms are the most important?
  - tf-idf (term frequency-inverse document frequency)
    - $\text{tfidf}(t,d,D) = \text{tf}(t,d) \times \text{idf}(t,D)$
  - $\text{tf}(t,d) = \text{length}(t|d) / \text{length}(d)$ 
    - frequency of term  $t$  in document  $d$ , divided by total number of terms in the document
  - $\text{idf}(t,D) = \log[ \text{length}(D) / \text{length}(D|t) ]$ 
    - inverse frequency of term  $t$  in corpus of documents  $D$

# How related two pieces of text are to each other

- you shall know a word by the company it keeps
  - LSA (latent semantic analysis)
  - PMI (pointwise mutual inclusion)
  - NGD (normalized google distance)
  - many others:
    - [https://en.wikipedia.org/wiki/Semantic\\_similarity](https://en.wikipedia.org/wiki/Semantic_similarity)

# PMI

- $\text{pmi}(x,y) = \log_2( p(x,y) / (p(x) \times p(y)) )$ 
  - log-scaled probability of  $x$  and  $y$  appearing together divided by probabilities of each
- if you just want to rank a set of terms  $y_1 \dots y_n$  in terms of how related they are to  $x$ , you don't need to know probabilities, just frequencies:
  - $\text{rank}(x,y_i) = f(x,y_i) / f(y_i)$



HELLO

BONJOUR

你好

07



# NATURAL LANGUAGE PROCESSING

# Assignment 8 (due before next lecture)

- Each team
  - is responsible for adding 2 questions based on Lectures 8 to the exam
- Team ambassador will
  - work with other teams to make sure your team's questions are unique
- Questions must be
  - multiple choice
  - with answer key
  - "Shuffle option order" (if it makes sense)
- Add your questions directly to shared form, AND paste your team's questions on blackboard

# Lab 8

- pick 5-7 different words
  - get PMI rank-scores for each pair of different words based on Google searches
- submit all word pairs with their PMI-rank scores, sorted by the scores on blackboard

# Free Parameters & Optimization

---

CS420 - Lectures 9-10  
Instructor: Dr. V



# Free Parameters

- Our RL agent had 3 parameters
  - learning rate
  - exploration parameter
  - discount factor for past or future states
- Our Neural Nets had a bunch of parameters too
  - learning rate
  - how many hidden layers
  - how many nodes in each hidden layer
  - which activation function to use (e.g., relu, step)

# Free model parameters

- Choosing correct values for free parameters is a common problem in AI
- How would you choose the values for each parameter, for example, in the Trace RL agent you created?
  - learning rate: .1? .01? .25?
  - exploration: .2? .05? .001?
  - decay: .9? .5? .1?

# Hill climbing

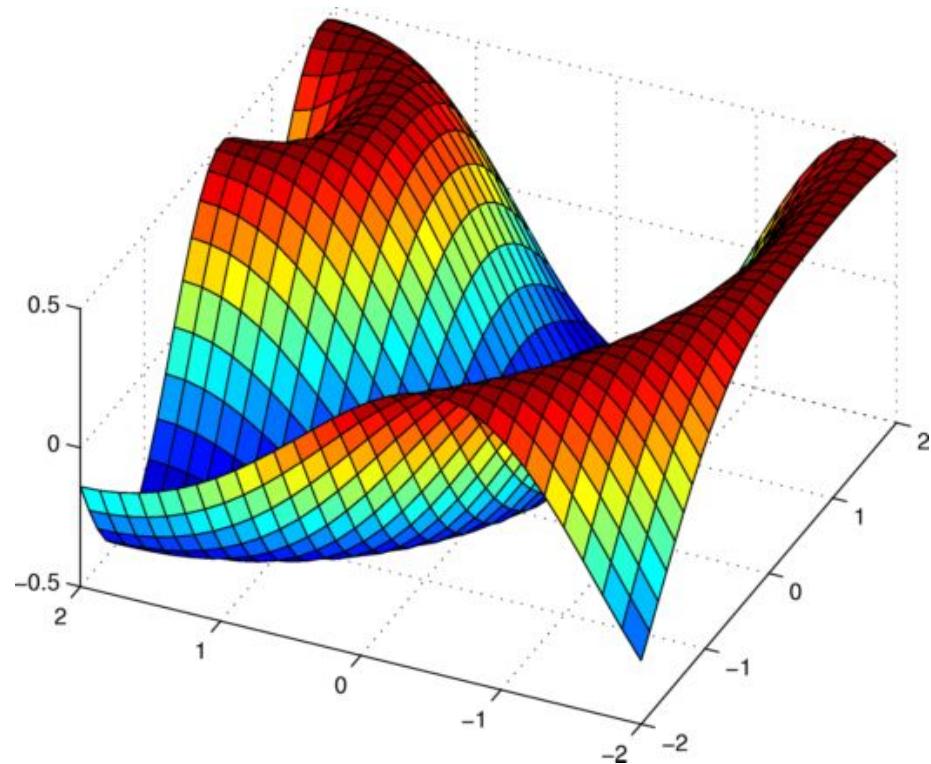
- make a slight change to a parameter value
  - if it results in better performance, keep going in the same direction, making slight changes
  - if it results in worse performance, go in a different direction
  - if value performance no longer improves, you found the best value for the parameter!
- repeat for every parameter

# Hillclimbing seems good...

- Can you foresee a problem with hill-climbing?
- What's another way to find [near-]optimal parameter values?

# Exhaustive search

- rather than doing anything smart, we can just look up **\*all\*** parameter value combinations
  - create a mesh grid
  - look up agent/model score for every point in grid



# grid search

```
import itertools, numpy as np

# 3 parameters, check all values from 0.05 to .95, by step of .05
vals = np.arange(.05,.95,.05)
for paramValues in itertools.product( vals,vals,vals ):
    score = getScore( *paramValues )

    ...

# if higher scores are better, report argmax, otherwise argmin
# i.e., parameter values that produce best score
```

# grid search

```
import itertools, numpy as np

# 3 parameters, check all values from 0.05 to .95, by step of .05
vals = np.arange(.05,.95,.05)
for paramValues in itertools.product( vals,vals,vals ):
    agent = TraceRL(ACTIONS, *paramValues)
    score = getScore( agent )
    ...
# if higher scores are better, report argmax, otherwise argmin;
# i.e., parameter values that produce best score
```

Full grid search seems good, but...

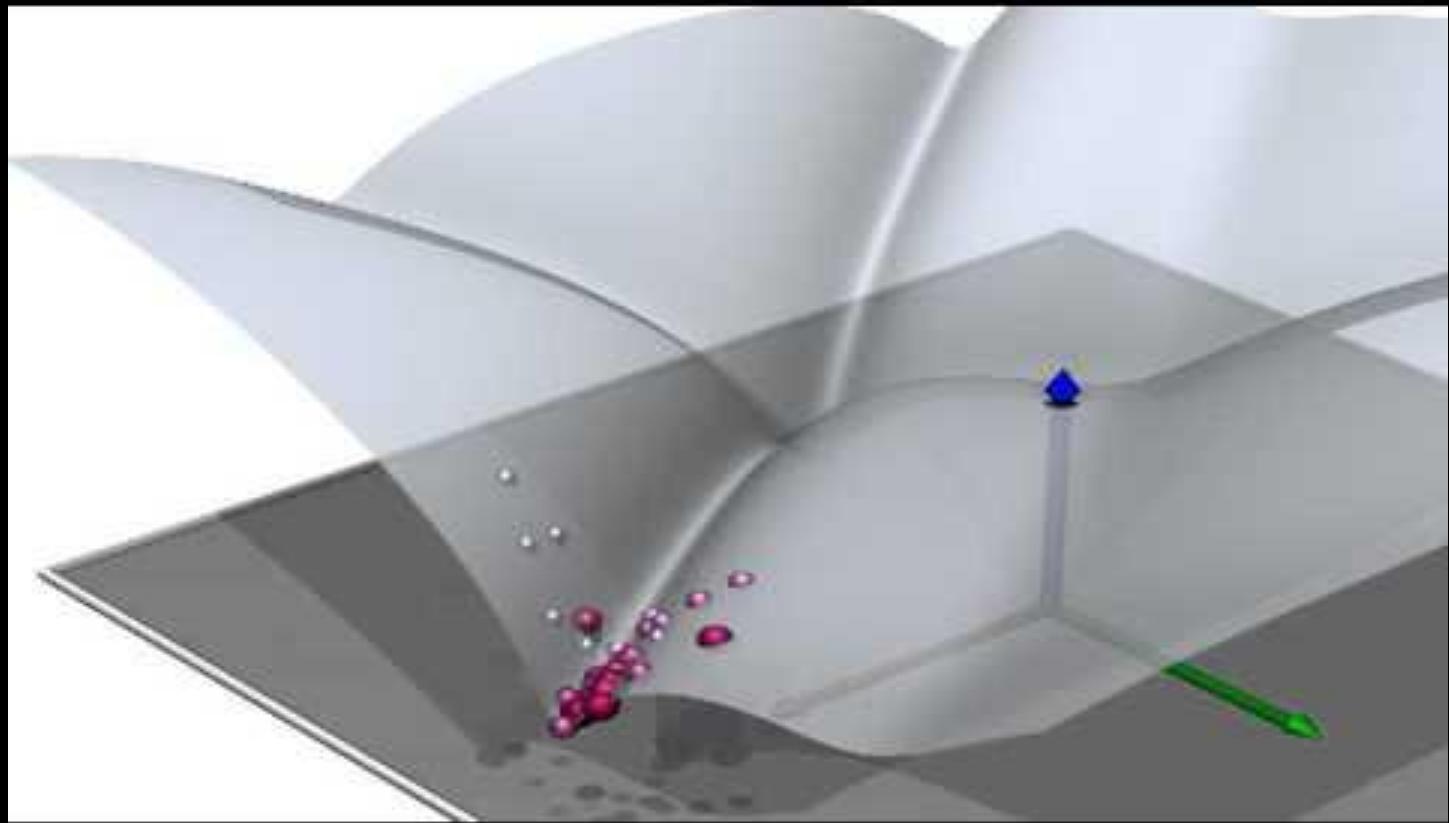
- Can you foresee a problem with an exhaustive grid-search?

# Full grid search seems good, but...

- Can you foresee a problem with an exhaustive grid-search?
  - if the step size is small, you'll find your near-optimal solution, but
    - time to solution increases exponentially with every additional free parameter
  - and if the step size is large, you'll find a solution much faster, but there will be gaps large enough to have missed better parameter values in your parameter space

# Simulated Annealing

- Simulated Annealing - a type of stochastic hill-climbing
- Simulated Annealing is like hill-climbing, but
  - there is a balance between exploration and exploitation
  - if you're still in exploratory mood (high temperature), you can move onto worse parameter values
  - as you get more greedy (lower temperatures), your tolerance for worse parameter values decreases



# Dual Annealing

- Dual Annealing is a version of Simulated Annealing that combines the global stochastic search with a second phase that includes a local search

## scipy.optimize.dual\_annealing

```
from scipy.optimize import dual_annealing
bounds = [(0.01,.5), (0.01,.5), (0.2,.99)]
result = dual_annealing(
    getScore,
    bounds)
# takes 80m on a high-end laptop; getScore runs in ~0.4s
print( result )
```

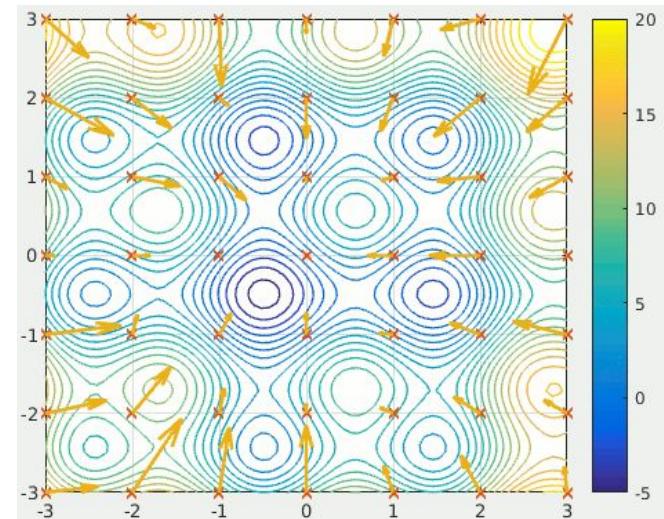
## scipy.optimize.dual\_annealing

```
from scipy.optimize import dual_annealing
bounds = [(0.01,.5), (0.01,.5), (0.2,.99)]
result = dual_annealing(
    lambda p:getScore(TraceRL(ACTIONS,*p)),
    bounds)
# takes 80m on a high-end laptop; getScore runs in ~0.4s

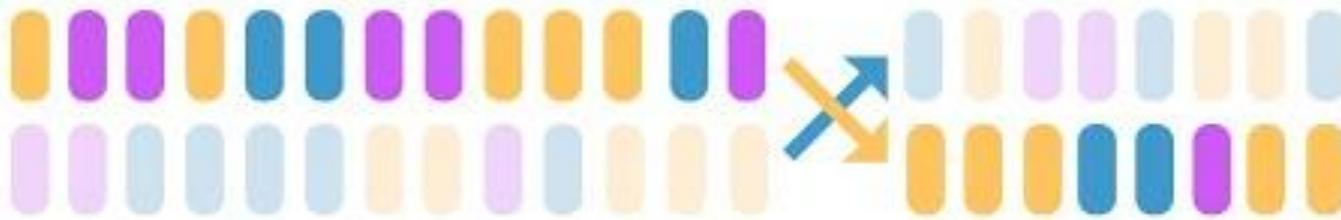
result
score: 526.72
message: ['Maximum number of iteration reached']
nfev: 7269
x: array([0.06, 0.09, 0.82])
```

# Optimization Algorithms

- There are many optimization algorithms
  - if a function is differentiable (you can find its derivative)
    - use stochastic gradient descent
  - if not
    - grid search
    - simulated annealing
    - evolutionary algorithms
    - particle swarm optimization ➔
    - cross-entropy method
    - etc



# GENETIC ALGORITHMS



EXPLAINED BY EXAMPLE

# Genetic Algorithms

- Simulates genetic evolution (natural selection)
  - first generation:
    - create random individual parameter sets (genes)
  - for every following generation:
    - **select fittest** individuals to reproduce
      - for every pair of parents, take some parameter values from each parent to produce a new set of **crossover** parameter values – a new genome
      - slightly **mutate** the new genome
    - keep the fittest individual(s) from current generation, delete the rest

## scipy.optimize.differential\_evolution

```
from scipy.optimize import differential_evolution
bounds = [(0.01,.5), (0.01,.5), (0.2,.99)]
result = differential_evolution(
    lambda p:getScore(TraceRL(ACTIONS,*p)),
    bounds, maxiter=50)
# takes 19.5m on a high-end laptop; getScore runs in ~0.4s

result
  score: 526.72
  message: ['Maximum number of iterations exceeded']
  nfev: 2371
convergence: 0.93987
  x: array([0.05, 0.02, 0.81])
```



P

NP

51  
3  
153

4	1	6	3	3	6	7	2	4
9	8	2	5	7	4	1	3	6
7	3	6	1	9	2	4	5	8
1	9	5	6	2	3	5	4	7
8	5	4	9	1	7	2	6	3
2	6	1	7	4	5	3	9	1
6	4	1	7	8	5	3	9	2
5	2	9	3	6	1	8	7	4
3	7	3	2	4	9	6	1	5

Sudoku



# Assignment 6 (due before next lecture)

- Each team
  - is responsible for adding 4 questions based on Lectures 9-10 to the exam
- Team ambassador will
  - work with other teams to make sure your team's questions are unique
- Questions must be
  - multiple choice
  - with answer key
  - "Shuffle option order" (if it makes sense)
- Add your questions directly to shared form, AND paste your team's questions on blackboard

# Lab 9

- use a grid search to find best values for
  - reinforcement learning agent parameters
- submit notebook on blackboard